

## Feature Selection / Extraction

The curse of dimensionality  
Feature extraction: Principal Component Analysis,  
Fisher Linear Discriminant Analysis  
Feature Selection: basis ideas

## Summary of episodes 1-3

- So far, we have seen that:
  - If class-conditional densities,  $p(x|\omega_i)$  and prior probabilities,  $P(\omega_i)$ , are known, one can set up optimal classifiers (e.g. likelihood ratio tests for dichotomy).
  - If only the parametric shape of class-conditional densities is known, it is possible to estimate their parameters using a set of labeled training samples.
  - If nothing is known, it is however possible to use non-parametric density estimation techniques to estimate class-conditional probabilities from a set of labeled training samples.
- In the next chapter, we will go further by considering the more difficult case when training data come unlabeled...
- Before that, in the present chapter, we focus on what happens when the dimensionality of the problem,  $D$ , becomes large

## The issue of dimensionality

- Intuitively, we might expect that considering more features helps increasing the discrimination between classes

- e.g. 2 Gaussian, equiprobable classes with same covariance:  $p(x|\omega_i) \sim \mathcal{N}(\mu_i, \Sigma)$ .  
It is easy to show that the probability of classification error is

$$P_e = \frac{1}{\sqrt{2\pi}} \int_{\theta=r}^{+\infty} \exp\left(-\frac{y^2}{2}\right) dy$$

where  $r$  is the Mahalanobis distance:  $r^2 = (\mu_1 - \mu_2)^T \Sigma^{-1} (\mu_1 - \mu_2)$

- In particular, if  $\Sigma = \text{diag}(\sigma_i^2)$

$$r^2 = \sum_{i=1}^D \left( \frac{\mu_{1i} - \mu_{2i}}{\sigma_i} \right)^2$$

- In other words, the larger  $D$ , the larger  $r$ , hence, the smaller  $P_e$  should be!

- This is not what can be observed in practice!

- Classification performance degrades as  $D$  becomes large...

## Trunk's experiment [Trunk, PAMI 1979]

Example where

$$\omega_1 \sim \mathcal{N}(\mu, I_D)$$

$$\omega_2 \sim \mathcal{N}(-\mu, I_D)$$

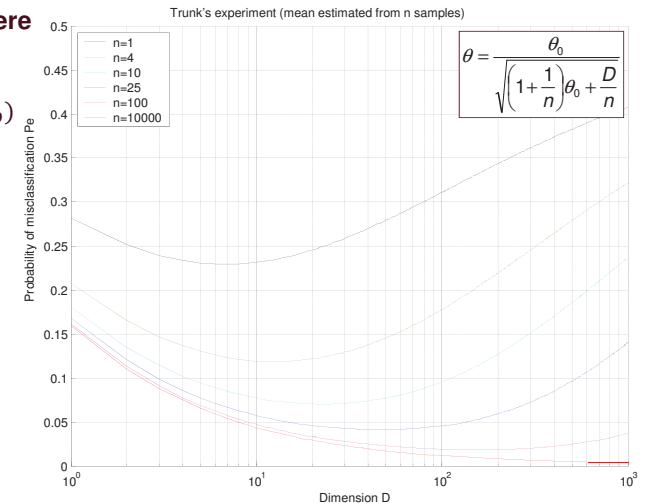
$$\mu_i = 1/\sqrt{i}$$

Then:

$$P_e = \int_{\theta}^{+\infty} \frac{e^{-y^2/2}}{\sqrt{2\pi}} dy$$

with  $\theta = r/2$

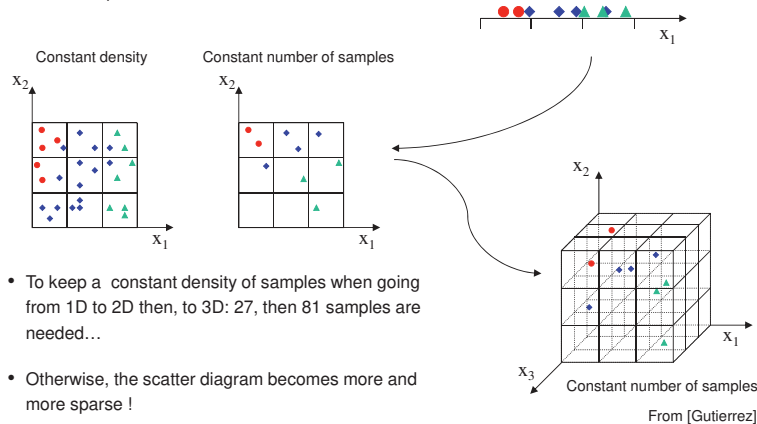
$$\theta_0 = \sum_{i=1}^{i=D} \frac{1}{i}$$



## The issue of dimensionality

### A theoretical study of the problem is very difficult [Duda]

- An explanation: the number of samples is finite, so PDF estimation is not very accurate...
- Example:



- To keep a constant density of samples when going from 1D to 2D then, to 3D: 27, then 81 samples are needed...
- Otherwise, the scatter diagram becomes more and more sparse !

## The curse of dimensionality ?

### Several problems occur as D increases

- Theoretical: scatter diagrams become more and more sparse, distances and volumes have weird properties in high dimension... [Bishop]
- Practical: mean, covariance matrices yield heavier computational burden...

### The set of problems that arise in data analysis when dimensionality increases is called the “curse of dimensionality” [Belmann 1961]

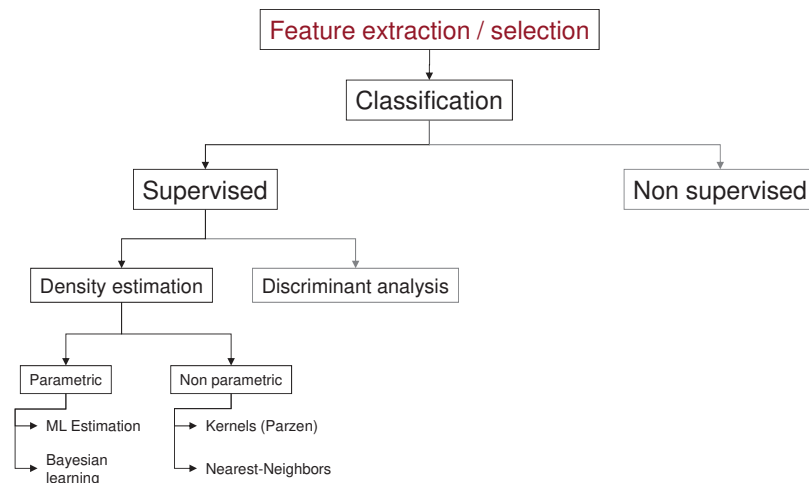
- There are several remedies (see e.g. [Duda])
- One of them is to apply a *dimensionality reduction* technique

### Another motivation for dimension reduction is visualization (needs $D \leq 3$ ...)

### Feature selection or extraction techniques

- Completely redesign the feature extraction algorithm...
- Select a subset of features or re-combine features


## A hierarchy of methods



## Dimensionality reduction techniques

### Two approaches:

- Combine features and select a reduced number of the resulting ones  
= Feature selection in transformed space or feature extraction
- Select features directly in feature space: feature selection




- We will study 2 well-known, linear feature extraction techniques: Principal Component Analysis (PCA) and Fisher's Linear Discriminant Analysis (LDA).
- Then, we will introduce nonlinear extraction techniques, and selection methods

## Principal Component Analysis (PCA)

- The choice of the transformation is generally guided by an optimization criterion

- In the case of PCA,

- The criterion assesses the quality of representation on a smaller-dimensional subspace
- Maximizing quality amounts to defining a rotation of the coordinate system



From [Gutierrez]

- A rotation of the coordinate system in feature space may reveal the structure of the data (see above “toy example”)
  - This is what PCA does, with a particular meaning of the word “structure”...

## Case 1: best representation by one point

- Goal: find the best representation of D-dimensional vectors,  $\mathbf{x}_k$ ,  $k=1..n$ , by a unique vector,  $\mathbf{x}_0$


- Representation quality criterion: mean squared error

$$J_0(\mathbf{x}_0) = \frac{1}{n} \sum_{k=1}^n \|\mathbf{x}_k - \mathbf{x}_0\|^2$$

- We seek  $\mathbf{x}_0$  that minimizes this criterion, i.e. the solution of:

$$\frac{\partial J_0(\mathbf{x}_0)}{\partial \mathbf{x}_0} = 0 = -\frac{2}{n} \sum_{k=1}^n (\mathbf{x}_k - \mathbf{x}_0) \longrightarrow \mathbf{x}_0 = \frac{1}{n} \sum_{k=1}^n \mathbf{x}_k = \boldsymbol{\mu}$$

- The sample mean is the best “0-dimensional” representation of the point cloud in the sense of mean square estimation



Ex : D = 2

## Case 2: best representation by a straight line

- Goal:

- Find the best 1D representation of the data, according to the model:

$$\mathbf{x} \approx \boldsymbol{\mu} + y \cdot \mathbf{e} \quad \text{with} \quad \|\mathbf{e}\|=1$$

- Straight line running through the sample mean,  $\boldsymbol{\mu}$  with (unknown) direction  $\mathbf{e}$ .
- Scalar,  $y$ : distance of any point  $\mathbf{x}$  lying on the line from  $\boldsymbol{\mu}$ .
- In particular, we have for any  $k$ :

$$\mathbf{x}_k \approx \boldsymbol{\mu} + y_k \cdot \mathbf{e}$$

- Mean squared error (MSE):  $J_1(y_1, \dots, y_n, \mathbf{e}) = \frac{1}{n} \sum_{k=1}^n \|\boldsymbol{\mu} + y_k \cdot \mathbf{e} - \mathbf{x}_k\|^2$

- Remark:

$$\begin{aligned} J_1(y_1, \dots, y_n, \mathbf{e}) &= \frac{1}{n} \sum_{k=1}^n \|y_k \cdot \mathbf{e} - (\mathbf{x}_k - \boldsymbol{\mu})\|^2 \\ &= \frac{1}{n} \sum_{k=1}^n \|y_k \cdot \mathbf{e}\|^2 - \frac{2}{n} \sum_{k=1}^n y_k \cdot \mathbf{e}^T (\mathbf{x}_k - \boldsymbol{\mu}) + \frac{1}{n} \sum_{k=1}^n \|\mathbf{x}_k - \boldsymbol{\mu}\|^2 \end{aligned}$$

## Representation by a line: the coordinates

- The set of coefficients  $y_k$  that minimize the squared error criterion is such that:

$$\frac{\partial J_1(y_1, \dots, y_n, \mathbf{e})}{\partial y_k} = 0 \quad \forall k = 1..n$$

- with  $J_1(y_1, \dots, y_n, \mathbf{e}) = \frac{1}{n} \sum_{k=1}^n y_k^2 - \frac{2}{n} \sum_{k=1}^n y_k \cdot \mathbf{e}^T (\mathbf{x}_k - \boldsymbol{\mu}) + \frac{1}{n} \sum_{k=1}^n \|\mathbf{x}_k - \boldsymbol{\mu}\|^2$


- which leads to  $0 = \frac{2}{n} y_k - \frac{2}{n} \mathbf{e}^T (\mathbf{x}_k - \boldsymbol{\mu})$

$$y_k = \mathbf{e}^T (\mathbf{x}_k - \boldsymbol{\mu})$$

- The coordinates on the line are given by the orthogonal projection of the (centered) data on  $\mathbf{e}$ .


## PCA vs. Orthogonal regression

$$y_k = \mathbf{e}^T (\mathbf{x}_k - \boldsymbol{\mu})$$



- PCA is akin to orthogonal regression

$$J_1(y_1, \dots, y_n, \mathbf{e}) = \frac{1}{n} \sum_{k=1}^n \|\boldsymbol{\mu} + y_k \mathbf{e} - \mathbf{x}_k\|^2$$



## Case 2: Re-formulating the problem

- Let us identify in  $J_1$  the expression of the newly calculated  $y_k$ 's

$$J_1(y_1, \dots, y_n, \mathbf{e}) = \frac{1}{n} \sum_{k=1}^n \underbrace{y_k^2 \|\mathbf{e}\|^2}_{\text{newly calculated } y_k} - \frac{2}{n} \sum_{k=1}^n \underbrace{y_k \mathbf{e}^T (\mathbf{x}_k - \boldsymbol{\mu})}_{\text{newly calculated } y_k} + \frac{1}{n} \sum_{k=1}^n \|\mathbf{x}_k - \boldsymbol{\mu}\|^2$$

- Thus, we obtain:

$$J_1(y_1, \dots, y_n, \mathbf{e}) = -\frac{1}{n} \sum_{k=1}^n y_k^2 + \frac{1}{n} \sum_{k=1}^n \|\mathbf{x}_k - \boldsymbol{\mu}\|^2$$

- Then:

$$J_1(y_1, \dots, y_n, \mathbf{e}) = -\frac{1}{n} \sum_{k=1}^n \mathbf{e}^T (\mathbf{x}_k - \boldsymbol{\mu}) (\mathbf{x}_k - \boldsymbol{\mu})^T \mathbf{e} + \frac{1}{n} \sum_{k=1}^n \|\mathbf{x}_k - \boldsymbol{\mu}\|^2$$

- And, finally :

$$J_1(y_1, \dots, y_n, \mathbf{e}) = -\frac{1}{n} \mathbf{e}^T \left( \sum_{k=1}^n (\mathbf{x}_k - \boldsymbol{\mu}) (\mathbf{x}_k - \boldsymbol{\mu})^T \right) \mathbf{e} + \frac{1}{n} \sum_{k=1}^n \|\mathbf{x}_k - \boldsymbol{\mu}\|^2$$

## Analyzing the new formulation of the problem

- We finally obtained two terms:

$$J_1(y_1, \dots, y_n, \mathbf{e}) = -\mathbf{e}^T \underbrace{\left( \frac{1}{n} \sum_{k=1}^n (\mathbf{x}_k - \boldsymbol{\mu}) (\mathbf{x}_k - \boldsymbol{\mu})^T \right)}_{\Sigma} \mathbf{e} + \underbrace{\frac{1}{n} \sum_{k=1}^n \|\mathbf{x}_k - \boldsymbol{\mu}\|^2}_{\text{independent from } \mathbf{e}}$$


- where  $\Sigma$  is the sample covariance matrix of the  $\mathbf{x}_k$ 's
- But, most of all,  $\mathbf{e}^T \Sigma \mathbf{e}$  is the variance of  $y$ , as shown below

- Note:  $y$  has zero mean (which can be straightforwardly checked!)


$$\text{Var}[y] = \frac{1}{n} \sum_{k=1}^n y_k^2 - \underbrace{\left[ \frac{1}{n} \sum_{k=1}^n y_k \right]^2}_0 = \frac{1}{n} \sum_{k=1}^n \mathbf{e}^T (\mathbf{x}_k - \boldsymbol{\mu}) (\mathbf{x}_k - \boldsymbol{\mu})^T \mathbf{e} = \mathbf{e}^T \Sigma \mathbf{e}$$

## Interpretation

- ⇒ Minimizing  $J_1$  with respect to  $\mathbf{e}$  ⇔ seeking the axis along which the variance (spread) of point coordinates is maximal...



☹ High projection error  
low variance



☺ Low projection error  
high variance

## Case 2: Solving the problem

### ■ To determine $\mathbf{e}$ , we must:

- Maximize  $\mathbf{e}^T \Sigma \mathbf{e}$
- Under the constraint:  $\|\mathbf{e}\|^2 = \mathbf{e}^T \mathbf{e} = 1$

### ■ Usual method: Lagrange multipliers

- see Optimization course
- Let us define:  $J'_1 = \mathbf{e}^T \Sigma \mathbf{e} - \lambda (\mathbf{e}^T \mathbf{e} - 1)$

where  $\lambda$  is a Lagrange Multiplier

$$\frac{\partial J'_1}{\partial \mathbf{e}} = 0 = 2\Sigma \mathbf{e} - 2\lambda \mathbf{e} \longrightarrow \Sigma \mathbf{e} = \lambda \mathbf{e}$$

### ■ Remember we are trying to maximize $\text{var}(y) = \mathbf{e}^T \Sigma \mathbf{e} = \lambda \mathbf{e}^T \mathbf{e} = \lambda$ ,

- Hence,  $\lambda$  is chosen as the largest eigenvalue of the covariance matrix,  $\Sigma$  and  $\mathbf{e}$  is the corresponding eigenvector

## Case 3: generalization to $d$ dimensions

### ■ In that case, the model is written as

$$\mathbf{x} \approx \boldsymbol{\mu} + \mathcal{E} \cdot \mathbf{y}$$

- where  $\mathbf{y}$  is a coordinate vector and  $\mathcal{E}$  is a matrix whose columns are the vectors  $\mathbf{e}_i$ .

$$J_n(\mathbf{y}_1, \dots, \mathbf{y}_n, \mathbf{e}_1, \dots, \mathbf{e}_d) = \frac{1}{n} \sum_{k=1}^n \left\| \boldsymbol{\mu} + \sum_{i=1}^d \mathbf{y}_{ki} \mathbf{e}_i - \mathbf{x}_k \right\|^2$$


- It can be shown that the minimum w.r.t.  $\mathbf{y}$  is obtained for the projection:

$$\mathbf{y} = \mathcal{E}^T (\mathbf{x}_k - \boldsymbol{\mu})$$

- And that the column vectors of  $\mathcal{E}$ , which minimize  $J_n$ , are given by the eigenvectors  $\Phi$  of the sample covariance matrix  $\Sigma$ :  $\Phi^T \Sigma \Phi = \Lambda = \text{diag}\{\lambda_i\}$
- They are orthogonal and represent the principal axes (of variation, see next slide).

## Interpretation

### ■ 2 phases:



### ■ PCA is a rotation because $\Phi$ is orthonormal ( $\Phi^T \Phi = \Phi \Phi^T = \mathbf{I}$ )

- This rotation de-correlates the data because  $\Sigma_y$  becomes diagonal

$$\Sigma_y = E[\mathbf{y}\mathbf{y}^T] = E[\Phi^T (\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T \Phi] = \Phi^T \Sigma \Phi = \Lambda$$

- The variance on each axis is:

$$E[\mathbf{y}_i^2] = E[\Phi_i^T (\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T \Phi_i] = \Phi_i^T \Sigma \Phi_i = \lambda_i$$

## Reconstruction and truncation

### ■ Sample mean and covariance matrix



$$\boldsymbol{\mu} = \frac{1}{n} \sum_{k=1}^n \mathbf{x}_k \quad \Sigma = \frac{1}{n} \sum_{k=1}^n (\mathbf{x}_k - \boldsymbol{\mu})(\mathbf{x}_k - \boldsymbol{\mu})^T$$

### ■ Diagonalization (rotation)

$$\Sigma = \Phi \Lambda \Phi^T \quad \Phi = \begin{bmatrix} \vdots & \vdots & \vdots & \vdots \\ \phi_1 & \dots & \phi_d & \dots & \phi_n & \dots & \phi_D \\ \vdots & \vdots & \vdots & \vdots \end{bmatrix}$$

$$\mathbf{y} = \Phi^T (\mathbf{x} - \boldsymbol{\mu}) \quad \Lambda = \text{diag}(\lambda_i)_{i=1 \dots D}$$


$$\mathbf{x} = \boldsymbol{\mu} + \Phi \mathbf{y}$$



### ■ Projection on eigenspace (with truncation)

$$\mathbf{y} = \Phi_d^T (\mathbf{x} - \boldsymbol{\mu}) \quad \Phi_d = \begin{bmatrix} \vdots & \vdots & \vdots \\ \phi_1 & \dots & \phi_d \\ \vdots & \vdots & \vdots \end{bmatrix}$$

$$\mathbf{x} \approx \boldsymbol{\mu} + \Phi_d \mathbf{y}$$




## Variances

- Eigenvectors are sorted by decreasing order of the corresponding eigenvalues (i.e. variances)

- PCA axes are ordered by decreasing order of variance
- Variability is most often concentrated on the first axes

- In general,  $d \ll D \rightarrow$  dimensionality reduction.

- The percentage of explained variance,  $\frac{\sum_1^d \lambda_i}{\sum_1^D \lambda_i}$  provides a principled way of selecting  $d$



- Variances along neglected dimension can be approximated optimally (in the sense of Maximum Likelihood) using their mean

$$\sigma^2 = \frac{1}{D-d} \sum_{i=d+1}^D \lambda_i \quad [\text{Moghaddam}]$$

- Note that  $\text{rank}(\Sigma) = \min(D, n-1) \rightarrow n-1$  eigenvalues computed if  $(n-1) < D$

## Probabilistic PCA (PPCA)

- Proposed by Tipping & Bishop (1997,1999) and Roweis (1998)
- Linear Generative Model

$$\mathbf{x} = \boldsymbol{\mu} + \mathcal{E}\mathbf{y} + \boldsymbol{\epsilon}$$

where:

- The latent variable,  $\mathbf{y}$ , is Gaussian:  $p(\mathbf{y}) \sim \mathcal{N}(\mathbf{0}, I)$
- The noise,  $\boldsymbol{\epsilon}$  is also Gaussian  $p(\boldsymbol{\epsilon}) \sim \mathcal{N}(\mathbf{0}, \sigma^2 I)$ , so:  $p(\mathbf{x}|\mathbf{y}) \sim \mathcal{N}(\boldsymbol{\mu} + \mathcal{E}\mathbf{y}, \sigma^2 I)$
- $\mathbf{y}$  and  $\boldsymbol{\epsilon}$  are independent.

- ML solution

$$\boldsymbol{\mu} = \frac{1}{n} \sum_{k=1}^n \mathbf{x}_k \quad \mathcal{E} \propto \Phi \quad \sigma^2 = \frac{1}{D-d} \sum_{i=d+1}^D \lambda_i$$

- PPCA brings several advantages compared with PCA, e.g.

- Derivation of an EM algorithm for fast computation of the first few eigenvectors
- PPCA model can be run generatively to draw samples from the distribution


## From PCA to LDA

- PCA

- Minimize a goodness-of-representation criterion
- Amounts to a rotation of coordinates in feature space
- ...in such a way as to represent most of the data variability on the first (few) dimensions of the transformed space.

- Fisher Discriminant Analysis (LDA)

- Optimizes a class-separability, or classification criterion.
- This does not provide the same directions
- But it remains a linear transform:  
 $\mathbf{y} = \mathbf{W}^T \mathbf{x}$



From [Gutierrez]

## Fisher Linear Discriminant Analysis (LDA)

- Case 1: two-class problem (dichotomy)

- Consider a set of  $n$  samples ( $D$ -dimensional vectors) :  $\{\mathbf{x}_1, \dots, \mathbf{x}_n\}$  among which  $n_1$  belong to class  $\omega_1$  and  $n_2$ , to class  $\omega_2$
- Projecting one sample on an axis  $\mathbf{w}$  yields the scalar  $y = \mathbf{w}^T \mathbf{x}$
- The vector  $\mathbf{w}$  is a projection direction (its norm does not really matter)

- We shall seek a direction  $\mathbf{w}$  along which the projections are as separable as possible

- In order to define a separability criterion, recall that

- The class means in feature space and in transformed space (for  $i=1,2$ ) are:

$$\mathbf{m}_i = \frac{1}{n_i} \sum_{\mathbf{x} \in \omega_i} \mathbf{x} \quad \tilde{\mathbf{m}}_i = \frac{1}{n_i} \sum_{y \in \omega_i} y = \frac{1}{n_i} \sum_{\mathbf{x} \in \omega_i} \mathbf{w}^T \mathbf{x} = \mathbf{w}^T \mathbf{m}_i$$

## LDA, case 1: two classes

- We should try to maximize

$$J(\mathbf{w}) = \tilde{m}_1 - \tilde{m}_2 = \mathbf{w}^T (\mathbf{m}_1 - \mathbf{m}_2)$$

$\rightarrow \mathbf{w} \propto (\mathbf{m}_1 - \mathbf{m}_2)$


- Does not account for intra-class variability
- To measure intra-class variability, one may use the variance (or the dispersion, which is proportional to the variance) along the axis,  $\mathbf{w}$ .

$$\tilde{s}_i^2 = \sum_{y \in \omega_i} (y - \tilde{m}_i)^2$$

- Fishers LDA maximizes:

$$J(\mathbf{w}) = \frac{(\tilde{m}_1 - \tilde{m}_2)^2}{\tilde{s}_1^2 + \tilde{s}_2^2}$$

$\left. \begin{array}{l} \text{• Separate means} \\ \text{• Keep classes compact} \end{array} \right\} \text{ along projection axis, } \mathbf{w}$



## LDA, case 1: two classes, expression of $J(\mathbf{w})$

- Goal: obtain  $J(\mathbf{w}) = \frac{(\tilde{m}_1 - \tilde{m}_2)^2}{\tilde{s}_1^2 + \tilde{s}_2^2}$  as a function of  $\mathbf{w}$

### 1. Denominator

$$\tilde{s}_i^2 = \sum_{y \in \omega_i} (y - \tilde{m}_i)^2 = \sum_{x \in \omega_i} (\mathbf{w}^T \mathbf{x} - \mathbf{w}^T \mathbf{m}_i)^2 = \mathbf{w}^T \sum_{x \in \omega_i} (\mathbf{x} - \mathbf{m}_i)(\mathbf{x} - \mathbf{m}_i)^T \mathbf{w} = \mathbf{w}^T S_i \mathbf{w}$$

$$\tilde{s}_1^2 + \tilde{s}_2^2 = \mathbf{w}^T (S_1 + S_2) \mathbf{w} = \mathbf{w}^T S_{\text{intra}} \mathbf{w}$$

### 2. Numerator

$$(\tilde{m}_1 - \tilde{m}_2)^2 = (\mathbf{w}^T \mathbf{m}_1 - \mathbf{w}^T \mathbf{m}_2)^2 = \mathbf{w}^T (\mathbf{m}_1 - \mathbf{m}_2)(\mathbf{m}_1 - \mathbf{m}_2)^T \mathbf{w} = \mathbf{w}^T S_{\text{inter}} \mathbf{w}$$

### 3. New expression

$$J(\mathbf{w}) = \frac{\mathbf{w}^T S_{\text{inter}} \mathbf{w}}{\mathbf{w}^T S_{\text{intra}} \mathbf{w}}$$

## LDA, case 1: maximizing $J(\mathbf{w})$

- Canceling the derivative of  $J(\mathbf{w})$  leads to

$$0 = \left[ \frac{\partial (\mathbf{w}^T S_{\text{inter}} \mathbf{w})}{\partial \mathbf{w}} \mathbf{w}^T S_{\text{intra}} \mathbf{w} - \mathbf{w}^T S_{\text{inter}} \mathbf{w} \frac{\partial (\mathbf{w}^T S_{\text{intra}} \mathbf{w})}{\partial \mathbf{w}} \right]$$

$$0 = \left[ \frac{\mathbf{w}^T S_{\text{intra}} \mathbf{w}}{\mathbf{w}^T S_{\text{intra}} \mathbf{w}} \right] S_{\text{inter}} \mathbf{w} - \left[ \frac{\mathbf{w}^T S_{\text{inter}} \mathbf{w}}{\mathbf{w}^T S_{\text{intra}} \mathbf{w}} \right] S_{\text{intra}} \mathbf{w} = S_{\text{inter}} \mathbf{w} - J(\mathbf{w}) S_{\text{intra}} \mathbf{w}$$

$$S_{\text{intra}}^{-1} S_{\text{inter}} \mathbf{w} = J(\mathbf{w}) \mathbf{w} \quad (\text{if } S_{\text{intra}} \text{ is non-singular})$$

- Remarks

- $S_{\text{inter}} \mathbf{w}$  is collinear to  $(\mathbf{m}_1 - \mathbf{m}_2)$ :  $S_{\text{inter}} \mathbf{w} = (\mathbf{m}_1 - \mathbf{m}_2) \underbrace{(\mathbf{m}_1 - \mathbf{m}_2)^T \mathbf{w}}_{\text{scalar}}$
- Since only the direction of  $\mathbf{w}$  is sought for  
→ the scale factor  $(\mathbf{m}_1 - \mathbf{m}_2)^T \mathbf{w} / J(\mathbf{w})$  can be neglected

$$\Rightarrow \hat{\mathbf{w}} = \arg \max_{\mathbf{w}} J(\mathbf{w}) = S_{\text{intra}}^{-1} (\mathbf{m}_1 - \mathbf{m}_2)$$

## Fisher Linear Discriminant (1936)

- The 1 dimensional projection direction which maximizes  $J$  is:

$$\hat{\mathbf{w}} = \arg \max_{\mathbf{w}} J(\mathbf{w}) = S_{\text{intra}}^{-1} (\mathbf{m}_1 - \mathbf{m}_2)$$

- This is not a discriminant, but a direction... So ???

- Once the direction of maximal separability has been found, it remains to determine a threshold  $\mathbf{x}_0$  along this axis. The classification rule is then to assign  $\mathbf{x}$  to class  $\omega_1$  if:

$$\hat{\mathbf{w}}^T (\mathbf{x} - \mathbf{x}_0) > 0$$

- This resembles what we have seen in the 2-class Gaussian case with equal covariance matrices! It can be shown that this is indeed the same solution...
- As a consequence, the above rule is optimal in the Gaussian case.


## Generalization to C classes (1/3)

- The generalization to C classes is quite straightforward

- In that case, we can find at most C-1 projection axes

- Tacitly, it is assumed that  $D \geq C$
- The vectors  $w_i$  are arranged in columns to form a matrix,  $W$
- The projections  $y_i$  are stored in a C-1 component vector,  $y$

$$y_i = w_i^T x \rightarrow y = W^T x$$



Three dimensional distributions are projected onto planes (described by their normal vectors  $w_1$  and  $w_2$ ). The best separability is in the plane associated with  $w_1$ .

Following [Duda]

## Generalization to C classes (2/3)

- General dispersion matrices

- Generalization of intra-class dispersion matrix :

$$S_{\text{intra}} = \sum_{i=1}^C S_i$$


$$\text{where } S_i = \sum_{x \in \omega_i} (x - m_i)(x - m_i)^T$$

$$\text{and } m_i = \frac{1}{N_i} \sum_{x \in \omega_i} x$$

- Generalization of inter-class dispersion matrix:

$$S_{\text{inter}} = \sum_{i=1}^C N_i (m_i - m)(m_i - m)^T$$

$$\text{where } m = \frac{1}{N} \sum_{x \in X} x = \frac{1}{N} \sum_{i=1}^C N_i m_i$$



$m$  is the center of mass of class means

Following [Gutierrez]

## Generalization to C classes (3/3)

- In transformed space, it can be shown that if

$$\tilde{S}_{\text{intra}} = \sum_{i=1}^C \sum_{y \in \omega_i} (y - \tilde{m}_i)(y - \tilde{m}_i)^T \quad \tilde{S}_{\text{inter}} = \sum_{i=1}^C N_i (\tilde{m}_i - \tilde{m})(\tilde{m}_i - \tilde{m})^T$$

$$\text{Then, } \tilde{S}_{\text{intra}} = W^T S_{\text{intra}} W \quad \text{and} \quad \tilde{S}_{\text{inter}} = W^T S_{\text{inter}} W$$

- We need a scalar function. To this end, we may use the determinant of dispersion matrices. Why ?

- The determinant of a (dispersion) matrix is the product of its eigenvalues
- i.e. product of variances in the principal directions. The determinant is thus homogeneous to the square of the hyperellipsoidal scattering volume (recall the gaussian case)

- The criterion can be written as:

$$J(W) = \frac{|\tilde{S}_{\text{inter}}|}{|\tilde{S}_{\text{intra}}|} = \frac{|W^T S_{\text{inter}} W|}{|W^T S_{\text{intra}} W|}$$

## Generalization to C classes, expression of $\hat{W}$

- The optimal projection axes are the solutions of a generalized eigenvector problem :

$$\hat{W} = [\hat{w}_1 | \hat{w}_2 | \dots | \hat{w}_{C-1}] = \arg\max_w \frac{|W^T S_{\text{inter}} W|}{|W^T S_{\text{intra}} W|} \Rightarrow (S_{\text{inter}} - \lambda_i S_{\text{intra}}) \hat{w}_i = 0$$

- Remarks:


- $S_{\text{inter}}$  is the sum of C matrices of rank 1 or 0 related by the definition of  $m$ . It is thus at most of rank C-1. This means that only C-1 eigenvalues are nonzero.
- LDA provides at most C-1 projection axes [Fukunaga90]
- They correspond to the largest eigenvalues of  $S_{\text{intra}}^{-1} S_{\text{inter}}$  (provided  $S_{\text{intra}}$  is nonsingular).
- The larger the eigenvalue, the higher the separability along the direction



## Linear extraction methods: summary

### PCA and LDA:

- Are widely used in practical applications
- Are not suited to complicated distributions (e.g. multi-modal, non Gaussian)...
- E.g. cases where LDA fails (following [Gutierrez])




### Other feature extraction techniques (quite numerous)

- Linear: Independent Component Analysis (ICA), Nonnegative Matrix Factorization (NMF), Multi-Dimensional Scaling (MDS)
- Non-linear: principal curves and surfaces, Kernel PCA, auto-encoders, ISOMAP, Linear Local Embedding (LLE), t-Distributed Stochastic Neighbor Embedding (t-SNE), Uniform Manifold Approximation and Projection (UMAP), etc. (a.k.a. manifold learning methods)

## Multi-Dimensional Scaling (MDS)

### Input: dissimilarity matrix

- Distance between samples in source space:  $\delta_{ij}$




### Output: positions

- low-dimensional target space ( $d=2$  or  $3$ ),  $d_{ij}$

### Property

- Inter-sample distances preserved in target space

E.g. MDS from 3D to 2D



### Optimization problem

- Minimize quadratic differences between  $\delta_{ij}$ 's and  $d_{ij}$ 's

### Unique up to global transformation

## ISOMAP

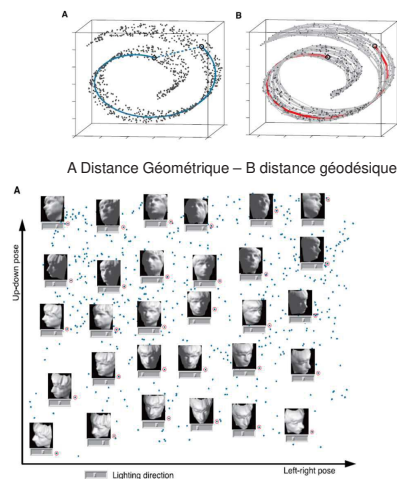
[Tenenbaum, Science, 2000]

### Embedding manifold

- May be non-linear


### ISOMAP =

- Similarity computation on manifold: geodesic distance
- Application of MDS



## Feature selection methods

### Feature Subset Selection



### Basis idea:

- Define an **objective criterion** to evaluate the optimality of a feature subset
- Identify such characteristics by optimizing this criterion using the training set.

### This seems rather simple !!! Unfortunately,

- Selecting  $t$  features among  $D$  involves

$$C_D^t = \binom{D}{t} = \frac{D!}{t!(D-t)!}$$

possibilities... E.g. 10 features among 25 = 3 268 760 possibilities !

⇒ Define a **research strategy**

## Feature selection methods

### Possible optimality criteria:

- Misclassification rate over a test set (must be different from training data!)
- Distance or separability measures:
  - ✓ Euclidean distance, Mahalanobis distance, etc...
  - ✓ Criteria based on dispersion matrices (e.g. LDA criterion)
  - ✓ Probabilistic (distance between probability densities)
- Correlation measures:
  - ✓ Features should be correlated with class labels, but not with each others.
  - ✓ Mutual Information measures between class labels and feature vectors.
- And more... (see [Webb, Duda, Theodoridis] for example).

### Optimization strategies

- Exhaustive search: optimal, but costly
- Tree search (e.g. "Branch and Bound" and variants)

## Branch and bound (1/2)


(following [Webb])

- Goal: **maximize a monotonic function** to find the optimal partition, avoiding to explore all possibilities !

$$X \subset Y \Rightarrow J(X) < J(Y)$$


- E.g. partition of 3 features among 5:  $C_5^3 = \frac{5!}{3!(5-3)!} = 10$  possibilities

- We may arrange possible partitions into a 10-leaves tree



## Branch and bound (2/2)

- Start from the rightmost leaf (note that the figure gives examples of J values) → J=77.2 becomes the current maximum
- Go back to the root and explore next branch → the current maximum becomes 80.1
- If we try to explore the next branch, we find J=60.9 at the first node. It is useless to continue because, by construction, the leaves will have smaller values. Same case for the last branch...
- In this example, we found the maximum in 8 evaluations of J instead of 10.



## Optimization strategies

### Sequential algorithms

- Sequential Forward Selection (SFS)
  - ✓ Starts from empty feature subset :  $Y_0 = \{\emptyset\}$
  - ✓ Sequentially adds the feature  $x^*$  which maximizes  $J(Y_k + x^*)$
  - ✓ Drawback: cannot discard features becoming obsolete due to the addition of other features...
- Sequential backward selection (SBS)
  - ✓ See SFS, but starting from the full set and discarding the feature that results in the smallest decrease of  $J(Y_k - x^*)$
- Plus L - minus R :
  - ✓ Combines SFS and SBS: at each step, discards R features and adds L features
  - ✓ Problem: no theoretical result for choosing L and R...

### Stochastic algorithms: escape from local minima !

- ✓ Random generation of a feature sample, then SFS, then SBS
- ✓ Simulated annealing, genetic algorithms, etc.

## Take-away remarks


- **To overcome the curse of dimensionality / ease visualization**
  - Feature selection or extraction
- **Selection:**
  - Given an optimality criterion (e.g. class separability)
  - Choice of features = combinatory optimization problem
- **Extraction**
  - Given a criterion that measures the:
    - ✓ Goodness of data representation (e.g. PCA)
    - ✓ Separability of classes (ex : LDA)
  - Optimization leads to a process of the form:
    - ✓ Linear transformation of data
    - ✓ Truncation
  - Other linear methods: ICA, NMF, MDS
  - Nonlinear methods, ISOMAP, LLE, auto-encoders, t-SNE, UMAP are better suited to complex variability.

## Exercise

### Feature selection and extraction

## Homework

- **Consider the following training sample:**
  - $X = (x_1, x_2) = \{ (1,2), (3,3), (3,5), (5,4), (5,6), (6,5), (8,7), (9,8) \}$
  - Draw the corresponding scatter plot
  - Calculate the sample mean and (biased) covariance matrix.
  - Calculate PCA and plot the eigenvectors  $v_1$  and  $v_2$ 
    - ✓ Trick: to calculate the first eigenvector easier, suppose that one of its coordinates is 1, calculate the other one, then normalize the resulting vector to 1.




From [Gutierrez]

## Solution of exercise 1

- $X = (x_1, x_2) = \{ (1,2), (3,3), (3,5), (5,4), (5,6), (6,5), (8,7), (9,8) \}$

$x_1$	$x_2$	$x_1 - \mu_1$	$x_2 - \mu_2$	$(x_1 - \mu_1)^2$	$(x_1 - \mu_1)(x_2 - \mu_2)$	$(x_2 - \mu_2)^2$
1	2	-4	-3	16	12	9
3	3	-2	-2	4	4	4
3	5	-2	0	4	0	0
5	4	0	-1	0	0	1
5	6	0	1	0	0	1
6	5	1	0	1	0	0
8	7	3	2	9	6	4
9	8	4	3	16	12	9
5	5			6,25	4,25	3,5

$$\Rightarrow \mu = \begin{pmatrix} 5 \\ 5 \end{pmatrix} \text{ and } \Sigma = \begin{bmatrix} 6,25 & 4,25 \\ 4,25 & 3,5 \end{bmatrix}$$



$$\begin{vmatrix} 6,25 - \lambda & 4,25 \\ 4,25 & 3,5 - \lambda \end{vmatrix} = 0 \Rightarrow \begin{cases} \lambda_1 \approx 9,34 \\ \lambda_2 \approx 0,41 \end{cases}$$

$$\begin{bmatrix} 6,25 & 4,25 \\ 4,25 & 3,5 \end{bmatrix} \begin{bmatrix} 1 \\ v \end{bmatrix} = 9,34 \cdot \begin{bmatrix} 1 \\ v \end{bmatrix} \xrightarrow{\text{(normalization)}} v_1 \approx \begin{pmatrix} 0,81 \\ 0,59 \end{pmatrix}, \text{ then } v_2 \perp v_1 \approx \begin{pmatrix} -0,59 \\ 0,81 \end{pmatrix}$$

[Gutierrez]



## Unsupervised learning and clustering

Parametric methods: mixture distributions – EM algorithm


Non-parametric methods: k-means algorithms, mode-seeking methods,

Flat vs. hierarchical algorithms

## Summary of previous episodes

- The classification methods that we have been studying so far are supervised (i.e. training samples are labeled)
  - When class-conditional distributions and prior probabilities are known, optimal classifiers can be designed in the Bayesian framework.
  - When only the functional forms of the pdf's are known, their parameters may be learned from a set of labeled training samples.
  - The pdf's may also be approximated in a non-parametric way from a set of labeled learning samples.
  - Dimensionality reduction techniques (feature selection or extraction) also use labeled samples.
- The situation is often far from being so ideal!
- What can be done when class labels are not known ?

## Unsupervised classification (a.k.a. clustering)



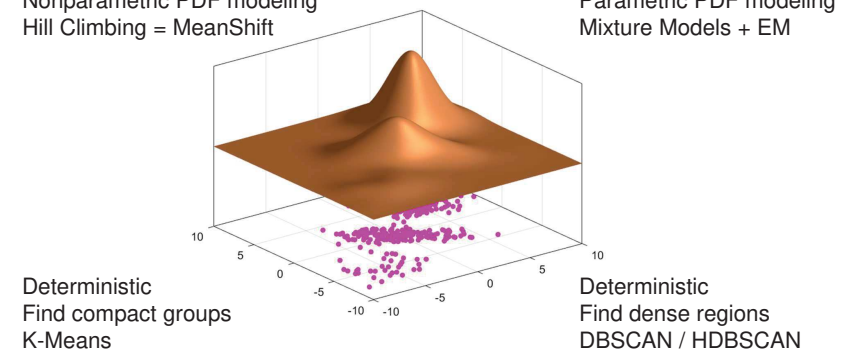
- Goal
  - Given a set of feature vectors  $X = \{x_1, \dots, x_N\}$  without class labels, capture the *structure* of the data
- Such a situation often occurs
  - Labeling a training set may be awkward.
  - Class labels are not always known in advance (e.g. *data mining*).
- NB: *Clustering* methods may also be used in a supervised context, for modeling multi-modal distributions.

## Clustering: deterministic vs. probabilistic

- “Hundreds of clustering algorithms have been proposed in the literature” [Jain2000]

Probabilistic  
Nonparametric PDF modeling  
Hill Climbing = MeanShift

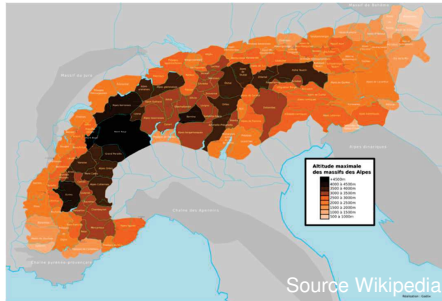
Probabilistic  
Parametric PDF modeling  
Mixture Models + EM



## Clustering: partitional vs. hierarchical

### Partitional (flat) approaches

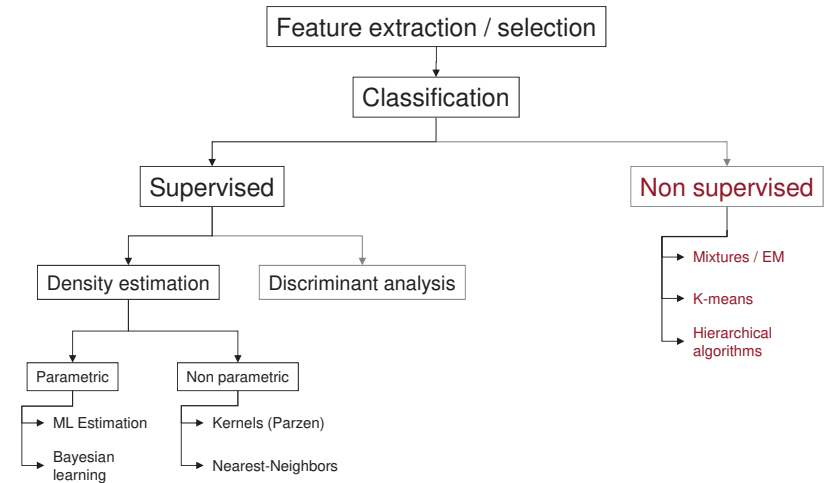
- Aim at partitioning feature space by optimizing a given functional



### Hierarchical approaches

- Organize the data into a series of nested groups that may be visualized as a tree-like representation

## A hierarchy of methods



## Parametric methods – Mixture Models

### They can model either

- A unique, complex (e.g. multimodal) probability density function
- A set of pdf's, each one being associated with one of the classes in the training set

### Assumptions


- The number of classes is supposed to be known
- The functional form of each class is known
- Each class may be of a different type
- Popular: **Gaussian Mixture Models (GMM)**

### The PDF of features is a mixture of parametric densities

$$p(\mathbf{x}|\Theta) = \sum_{j=1}^C p(\mathbf{x}|\omega_j, \theta_j) \pi_j$$

- Note: such approaches are sometimes dubbed *semi-parametric*.

## PDF estimator families



## Mixture models: the unknown parameters

- The mixing parameters,  $\pi_j$ , can be viewed as prior class probabilities

$$p(\mathbf{x}|\boldsymbol{\theta}) = \sum_{j=1}^C p(\mathbf{x}|\omega_j, \boldsymbol{\theta}_j) P(\omega_j)$$

- The vector of unknown parameters,  $\boldsymbol{\theta}$ , that concatenates:
  - The parameters of each component  $\boldsymbol{\theta}_j$  (e.g. for a Gaussian mixture:  $\boldsymbol{\mu}_j$  and  $\boldsymbol{\Sigma}_j$ )
  - The mixing parameters  $\pi_j = P(\omega_j)$  with  $\sum_{j=1}^C P(\omega_j) = 1$

must be estimated from unlabeled samples  $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$

- We first form the likelihood:

$$p(\mathbf{X}|\boldsymbol{\theta}) = \prod_{k=1}^N \sum_{j=1}^C p(\mathbf{x}_k|\omega_j, \boldsymbol{\theta}_j) P(\omega_j)$$

## Mixture models: Maximum Likelihood estimation

- It may be shown (see [DUDA], pp. 518-528) that the maximum likelihood solution:

$$L(\boldsymbol{\theta}) = \ln[p(\mathbf{X}|\boldsymbol{\theta})] = \sum_{k=1}^N \ln \left[ \sum_{j=1}^C p(\mathbf{x}_k|\omega_j, \boldsymbol{\theta}_j) P(\omega_j) \right]$$

satisfies in the Gaussian case,  $\boldsymbol{\Sigma}_j = \sigma_j^2 \mathbf{I}_d$  (in  $d$  dimensions):

$$\hat{\boldsymbol{\mu}}_j = \frac{\sum_k P(\omega_j|\mathbf{x}_k, \boldsymbol{\theta}) \mathbf{x}_k}{\sum_k P(\omega_j|\mathbf{x}_k, \boldsymbol{\theta})} \quad \hat{\sigma}_j^2 = \frac{\sum_k P(\omega_j|\mathbf{x}_k, \boldsymbol{\theta}) \|\mathbf{x}_k - \hat{\boldsymbol{\mu}}_j\|^2}{d \sum_k P(\omega_j|\mathbf{x}_k, \boldsymbol{\theta})}$$

for mixing parameters (under unit sum condition  $\sum_j \hat{P}(\omega_j) = 1$ )

$$\hat{P}(\omega_j) = \frac{1}{N} \sum_k P(\omega_j|\mathbf{x}_k, \boldsymbol{\theta})$$

## Interpretation

- The weights:  $w_{kj} = P(\omega_j|\mathbf{x}_k, \boldsymbol{\theta})$  can be seen as an estimation of the chance that  $\mathbf{x}_k$  belongs to the  $j$ -th class.

• NB:  $\sum_j P(\omega_j|\mathbf{x}_k, \boldsymbol{\theta}) = 1$

- Should  $P(\omega_j|\mathbf{x}_k, \boldsymbol{\theta})$  be binary ( $\Leftrightarrow$  one-hot encoding), then

$$\sum_k P(\omega_j|\mathbf{x}_k, \boldsymbol{\theta}) = \sum_{\mathbf{x}_k \in \omega_j} 1 = N_j \quad \sum_k P(\omega_j|\mathbf{x}_k, \boldsymbol{\theta}) \mathbf{x}_k = \sum_{\mathbf{x}_k \in \omega_j} \mathbf{x}_k$$

- The ML equations would become:

$$\hat{\boldsymbol{\mu}}_j = \frac{1}{N_j} \sum_{\mathbf{x}_k \in \omega_j} \mathbf{x}_k \quad \hat{\sigma}_j^2 = \frac{1}{N_j d} \sum_{\mathbf{x}_k \in \omega_j} \|\mathbf{x}_k - \hat{\boldsymbol{\mu}}_j\|^2 \quad \hat{P}(\omega_j) = \frac{N_j}{N}$$

- And we would get the usual estimators!

- Most of the time, however,  $P(\omega_j|\mathbf{x}_k, \boldsymbol{\theta}) \in [0,1]$

- All samples play a role in the estimation of all class parameters...

## Maximum Likelihood estimation: difficulty

- Moreover, in the ML equations:

$$P(\omega_j|\mathbf{x}_k, \boldsymbol{\theta}) = \frac{\overbrace{p(\mathbf{x}_k|\omega_j, \boldsymbol{\theta}) \hat{P}(\omega_j)}^{\frac{1}{(2\pi\hat{\sigma}_j^2)^{d/2}} \exp\left(-\frac{\|\mathbf{x}_k - \hat{\boldsymbol{\mu}}_j\|^2}{2\hat{\sigma}_j^2}\right)}}{\sum_c p(\mathbf{x}_k|\omega_c, \boldsymbol{\theta}) \hat{P}(\omega_c)}$$



- The ML equations are coupled... So?

- Idea: use a fixed-point algorithm

- Take an initial parameter set
- Use at each iteration the ML equations, inserting in their rightmost term the values found at the previous iteration...

- The EM framework is a principled way of devising such an algorithm, and of proving its convergence towards a (local) maximum of the log-likelihood.

## The EM algorithm

[Dempster77]

- EM means Expectation - Maximization
- Iterative algorithm for parameter estimation, in the sense of Maximum Likelihood, from missing or hidden data.
- Missing (or corrupted) data:
  - When the observation process is limited or corrupted.
- Hidden data:
  - Their knowledge would lead to a simpler form of likelihood
- Applications:
  - Tomographic reconstruction, mixture models...

## The EM formalism

- **Notations:**
  - $X$  denotes the incomplete, observed data
  - $Z$  denotes the missing, unknown data.
  - $(X, Z)$  are the complete data
  - $\Theta$ , is the parameter vector, which concatenates  $\theta$ , the parameters of the components and  $P(\omega)$ , the mixing parameters.
- We may form the likelihood of the complete data:  $p(X, Z | \Theta)$
- In contrast,  $p(X | \Theta)$ , is called likelihood of the incomplete data
- Maximizing  $L(\Theta) = \log[p(X | \Theta)]$  is difficult...
  - Maximizing  $\log[p(X, Z | \Theta)]$  may be easier, for a particular choice of  $Z$ .
  - But  $Z$  is unobservable...  $\rightarrow$  we shall try to eliminate it!

## EM: basis ideas (1/2)

- Link between the likelihood of the complete data and the likelihood of the incomplete data
  - According to the definition of conditional probabilities, we have:

$$p(X | \Theta) = p(X, Z | \Theta) / p(Z | X, \Theta)$$

- Taking logs, we have:

$$L(\Theta) = \log[p(X | \Theta)] = \log[p(X, Z | \Theta)] - \log[p(Z | X, \Theta)]$$

- 1<sup>st</sup> idea: if we average this expression with respect to  $Z$ ,
  - The leftmost term (log-likelihood) does not change...
  - And we eliminate  $Z$  from the rightmost term!

## EM: basis ideas (2/2)

- In an iterative process, we know the current estimate,  $\Theta^n$ 
  - The average (expectation) is taken with respect to  $Z$ , knowing  $X$  and  $\Theta^n$
- We obtain:

$$L(\Theta) = E_Z[\log p(X, Z | \Theta) | X, \Theta^n] - E_Z[\log p(Z | X, \Theta) | X, \Theta^n]$$

...which may be written as:

$$L(\Theta) = \log[p(X | \Theta)] \equiv Q(\Theta | \Theta^n) - H(\Theta | \Theta^n)$$

- 2<sup>nd</sup> idea: take advantage of a “good property” of the expectation to get rid of  $H$ !



## A “good property” of the expectation

- It can be shown [Lange2000] that

$$H(\Theta^n | \Theta^n) - H(\Theta^{n+1} | \Theta^n) \geq 0$$



- In the EM,  $\Theta^{n+1}$  is chosen to maximize  $Q(\Theta | \Theta^n)$

- So, in particular:

$$Q(\Theta^{n+1} | \Theta^n) - Q(\Theta^n | \Theta^n) \geq 0$$

- Summing up:

$$\begin{aligned} & Q(\Theta^{n+1} | \Theta^n) - Q(\Theta^n | \Theta^n) \geq 0 \\ + & -H(\Theta^{n+1} | \Theta^n) - \{-H(\Theta^n | \Theta^n)\} \geq 0 \\ = & L(\Theta^{n+1}) - L(\Theta^n) \geq 0 \end{aligned}$$

→ which shows that each iteration increases the log-likelihood

## The iteration of the EM algorithm

- The (n+1)-th iteration is made of 2 steps

- E (expectation) step:** form the conditional expectation

$$Q(\Theta | \Theta^n) = E_z[\log p(X, Z | \Theta) | X, \Theta^n] = \int \log p(X, z | \Theta) p(z | X, \Theta^n) dz$$


- M (maximization) step:** find  $\Theta$ , the maximizer of  $Q(\Theta | \Theta^n)$

$$\Theta^{n+1} = \arg \max_{\Theta} Q(\Theta | \Theta^n)$$

- Since the M-step always increases  $L(\Theta) = \log[p(x|\Theta)]$ , the algorithm converges towards a stationary point (which may not be the global maximum!)
- In many cases, the E-step is performed analytically and grouped with the M-step (the iteration is implemented as a single closed-form expression)

## The EM algorithm: interpretation

- During the E-step, we eliminate the unobserved data  $Z$  by integration (marginalization)
- During the M-step, we maximize the average likelihood of the complete data



From [Gutierrez]

## The case of mixture models

- The hidden variable  $Z = \{z_k\}$  for  $k=1 \dots N$ , indicates which component (class)  $x_k$  has been generated from
  - Each  $z_k$  is a vector with  $C$  components,  $z_{kj}$
  - $z_{kj}$  is 1 if  $x_k$  belongs to class  $\omega_j$  and 0 otherwise
  - $z_{kj}$  is a Bernoulli variable with parameter  $P$ =probability that  $x_k$  belongs to class  $\omega_j$  (knowing the parameters), i.e.  $P(\omega_j | x_k, \Theta)$
  - Since the expectation of a Bernoulli law with parameter  $P$  is ...  $P$ , we have

$$E(z_{kj} | x_k, \Theta) = P(\omega_j | x_k, \Theta) \stackrel{\text{def}}{=} w_{kj}$$

- This expression may be evaluated using Bayes' theorem

$$w_{kj} = P(\omega_j | x_k, \Theta) = \frac{P(x_k | \omega_j, \Theta) P(\omega_j)}{\sum_c P(x_k | \omega_c, \Theta) P(\omega_c)}$$

- Let us note these results and form the complete likelihood...

## Mixture models: the E-step

- The likelihood of the complete data  $(x_k, z_k)$  is written as

$$p(x_k, z_k | \Theta) = p(x_k | z_k, \Theta) P(z_k | \Theta) = p(x_k | \theta_j) P(\omega_j) = \prod_{j=1}^C [p(x_k | \theta_j) P(\omega_j)]^{z_{kj}}$$

- For N independent observations:

$$p(X, Z | \Theta) = \prod_{k=1}^N \prod_{j=1}^C [p(x_k | \theta_j) P(\omega_j)]^{z_{kj}}$$

$$\log(p(X, Z | \Theta)) = \sum_{k=1}^N \sum_{j=1}^C z_{kj} [\log(p(x_k | \theta_j)) + \log(P(\omega_j))]$$

- Now, we take the expectation with respect to  $z$

- Recalling the results of the previous slide, it comes that:

$$Q(\Theta | \Theta^n) = \sum_{j=1}^C \sum_{k=1}^N \log(P(\omega_j)) w_{kj}^n + \sum_{j=1}^C \sum_{k=1}^N \log(p(x_k | \theta_j)) w_{kj}^n$$

$$w_{kj}^n = P(\omega_j | x_k, \Theta^n)$$

## Gaussian mixtures: the M-step

- Canceling the derivative of Q, with respect to each parameter, one obtains in the Gaussian case

$$w_{kj}^n = P(\omega_j | x_k, \Theta^n)$$

$$= \frac{p(x_k | \omega_j, \theta^n) P(\omega_j)}{\sum_c p(x_k | \omega_c, \theta^n) P(\omega_c)}$$

$$\frac{1}{(2\pi)^{d/2} |\Sigma_j^n|^{1/2}} \exp\left(-\frac{(x_k - \mu_j^n)(\Sigma_j^n)^{-1}(x_k - \mu_j^n)}{2}\right)$$

$$\mu_j^{n+1} = \frac{\sum_k w_{kj}^n x_k}{\sum_k w_{kj}^n}$$

$$\Sigma_j^{n+1} = \frac{\sum_k w_{kj}^n (x_k - \mu_j^{n+1})(x_k - \mu_j^{n+1})^T}{\sum_k w_{kj}^n}$$

$$P^{n+1}(\omega_j) = \frac{1}{N} \sum_k w_{kj}^n$$

- Note that

- We obtain weighted expressions for the mean and covariance matrix
- The implementation is rather simple, but the convergence is slow
- The closer  $x_k$  from the class mean, the higher its weight

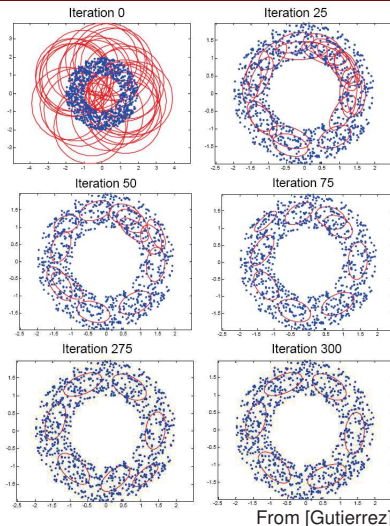


## Example (modeling a single, complex pdf)

- Training data:  
900 points drawn from an annular distribution


- EM algorithm:

- 30 initial classes with random means
- Diagonal initial covariance matrices with large variance
- Regularization of covariance matrices
- Components with small  $P(\omega_i)$  are discarded
- The final model is compact : only a few components



## Unsupervised classification for 3 classes

- Simulated data: mixture of 3 bidimensional Gaussian classes.
- Initialization (left), solution after 50 iterations (right)



## EM clustering and classification

### Classification method

- Train:
  - Specify the number of classes and the parametric form of their pdf's
  - Estimate the parameters of the mixture from training data (using EM)
- Classify new data,  $\mathbf{x}$ :
  - Form posterior probabilities  $P(\omega_j|\mathbf{x})$  and chose the class with highest probability.

### Issues


- Choosing C
  - Minimizing a criterion with respect to C : log-likelihood + penalty term (penalize complex models i.e. large C values),  $k=\text{nb parameters}$ 
    - Akaike Information Criterion:  $AIC = 2L(\hat{\theta}) + 2k$
    - Bayesian Information Criterion:  $BIC = 2L(\hat{\theta}) + k \ln N$
  - See **EMbic.m** demo
  - Cross-validation, etc.
- The EM may be trapped in local maxima, and it is rather slow
  - Variants: GEM (Generalized EM), CEM (Classification EM), SEM (Stochastic EM)



## Mode seeking methods

### Assumptions

- In the overall PDF  $p(\mathbf{x})$ , each cluster corresponds to a mode.
- 2D example: "hills" separated by "valleys"



### Principle

- Seek the modes using e.g. gradient ascent algorithm
- Example: *Meanshift* [Comaniciu99] ...after [Fukunaga 75] !
  - $p(\mathbf{x})$  is modeled using Parzen windows :  $p_K(\mathbf{x}) \propto \sum_i k(\|\mathbf{x} - \mathbf{x}_i\|/h)^2$
  - Meanshift* is a hill-climbing method adapted to such a representation

### Gradient ascent and the Meanshift

- $\nabla p_K \propto \sum_i (\mathbf{x} - \mathbf{x}_i) k'(\|\mathbf{x} - \mathbf{x}_i\|/h)^2$
- The (normalized) gradient direction, provided that  $g(\mathbf{x}) = -k'(\mathbf{x})$ , is given by the meanshift


$$\frac{\nabla p_K}{p_K} \propto \frac{\sum_i x_i g\left(\left\|\frac{\mathbf{x} - \mathbf{x}_i}{h}\right\|^2\right)}{\sum_i g\left(\left\|\frac{\mathbf{x} - \mathbf{x}_i}{h}\right\|^2\right)} - \mathbf{x}$$

## The Meanshift

### Gradient ascent: $\mathbf{x} \leftarrow \mathbf{x} + \nabla p_K / P_G$ ,

### Algorithm

- Starting from each data point, iterate :  $\mathbf{x} \leftarrow \frac{\sum_i x_i g\left(\left\|\frac{\mathbf{x} - \mathbf{x}_i}{h}\right\|^2\right)}{\sum_i g\left(\left\|\frac{\mathbf{x} - \mathbf{x}_i}{h}\right\|^2\right)}$  until convergence
- Group points that converge to the same maximum into a cluster



[Comaniciu2002]

## Deterministic methods / K-means clustering

### Does not involve pdf estimation

- Attempts to find natural grouping (clusters) in a data set by optimizing an objective function
- Find an assignment of data points to clusters, and a set of prototypes that represent the centers of the clusters, in order to minimize

$$J_e = \sum_{j=1}^C \sum_{\mathbf{x} \in \omega_j} \|\mathbf{x} - \mathbf{m}_j\|^2$$


## Remark

### Dispersion matrices (see FDA)

$$S_{\text{intra}} = \sum_{j=1}^C \sum_{x \in \omega_j} (x - m_j)(x - m_j)^T$$

$$S_{\text{inter}} = \sum_{j=1}^C N_j (m_j - m)(m_j - m)^T$$

$$S_{\text{total}} = \sum_{j=1}^C \sum_{x \in \omega_j} (x - m)(x - m)^T = S_{\text{intra}} + S_{\text{inter}}$$



### Taking the trace, we obtain (Huygens theorem)

$$J_{\text{total}} = \sum_{j=1}^C \sum_{x \in \omega_j} \|x - m\|^2 = \sum_{j=1}^C \sum_{x \in \omega_j} \|x - m_j\|^2 + \sum_{j=1}^C N_j \|m_j - m\|^2 = J_e + J_{\text{inter}}$$

- Since the total inertia,  $J_{\text{total}}$  is constant for a given data set and a constant  $C$ , minimizing  $J_e$  amounts to maximizing the inter-class variance  $J_{\text{inter}}$

→ Finding compact clusters ⇔ finding distinct clusters

## K-means: basis idea [Lloyd 1957-1982, McQueen 1967]

### Let us introduce binary variables to formalize the assignment of data points to clusters (one-hot-encoding)

$$u_{kj} \in \{0,1\} \text{ and } \sum_{j=1}^C u_{kj} = 1$$

$$J_e(u, m) = \sum_{j=1}^C \sum_{k=1}^N u_{kj} \|x_k - m_j\|^2$$

### Fix the $m_j$ 's and minimize with respect to $u$

- For each  $x_k$ , only one  $u_{kj}$  may be non null (due to unit sum condition).
- Clearly, in order to minimize  $J_e$ , we shall choose (for each  $k$ ):

$$u_{kj} = \begin{cases} 1 & \text{if } \|x_k - m_j\|^2 < \|x_k - m_i\|^2 \forall i \\ 0 & \text{otherwise} \end{cases}$$

## K-means: basis idea [Lloyd 1957-1982, McQueen 1967]

### Assignment rule = nearest-mean rule

- Optimal for Gaussian clusters with equal prior and same, isotropic, covariance matrix

### Now, fix $u$ and minimize w.r.t. the $m_j$ 's

$$J_e(u, m) = \sum_{j=1}^C \sum_{k=1}^N u_{kj} \|x_k - m_j\|^2$$

$$0 = \sum_{k=1}^N u_{kj} \frac{\partial \|x_k - m_j\|^2}{\partial m_j} \rightarrow \sum_{k=1}^N u_{kj} x_k = \sum_{k=1}^N u_{kj} m_j = N_j m_j \rightarrow m_j = \frac{1}{N_j} \sum_{x \in \omega_j} x$$

→ Optimal class centers are class means

## The k-means algorithm

### Batch k-means algorithm

- Initialize class centers
  - ✓ At random, or simply take the first  $C$  samples, for example...
- Repeat
  - ✓ Classify the  $n$  samples according to their distance to class centers
    - ⇒ if  $\|x_k - m_j\| < \|x_k - m_i\| \forall i \neq j$ , then  $x_k$  is assigned to  $\omega_j$
  - ✓ Update the positions of class centers
    - ⇒  $m_j = \frac{1}{N_j} \sum_{x \in \omega_j} x \quad \forall j = 1 \dots C$
- Until the partition does not change



## Remarks (1)

- **The K-means algorithm is rather fast  $O(NdCT)$** 
  - where  $N$  is the number of samples,  $d$  the dimensionality of feature space,  $C$  the number of classes and  $T$  the number of iterations (in general,  $T \ll N$ ).
- **The K-means algorithm converges to a local minimum of  $J_e$ , which depends on the initial choice of class centers.**
  - One may run the algorithm several times, with different initializations and retain the best result (e.g. the one that yields the smallest  $J_e$  final value)
- **Dynamic clouds generalize k-means. They involve:**
  - A function that allows computing the prototype for each cluster
  - A function that allows assigning each point to one class
  - The general form of the algorithm remains the same...

## Remarks (2)

- **There exists a sequential version of the k-means algorithm [Duda, p. 549]**
  - Initialize class centers
  - Repeat
    - ✓ Pick one sample at random,  $x^*$
    - ✓ Assign  $x^*$  to a new class according to its distance to class centers
    - ✓ Update the positions of the centers of the old and new class for  $x^*$  and the value of the overall error,  $J_e$
  - Until  $J_e$  does not change in  $n$  attempts
- **This version is (even) less robust with respect to local minima...**

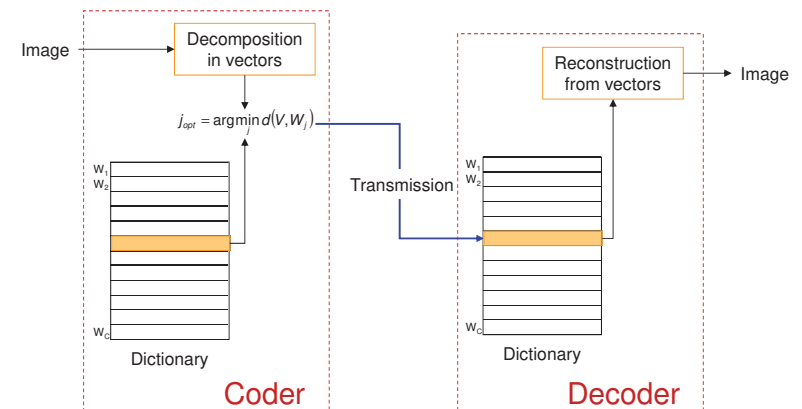
## The ISODATA algorithm

[Ball & Hall, 1965]

- **ISODATA is a more sophisticated version of k-means**
  - K-means are sometimes called "basic ISODATA"
  - Supplementary mechanisms allow the automatic selection of the number of clusters,  $C$
- **Algorithm:**
  - Repeat
    - ✓ Perform clustering using k-means
    - ✓ Discard clusters smaller than a certain threshold
    - ✓ Split highly dispersed clusters
    - ✓ Merge sufficiently close clusters
  - Until stabilization
- **Thresholds must be set, but the algorithm is more automatic**

## K-means and vector quantization

- **Vector quantization**




## K-means and vector quantization

### Construction of the dictionary: minimization of total distortion


- Most often, distortion = quadratic error (i.e. k-means criterion)

### This is a clustering problem!

- A famous algorithm (called LBG) was proposed in [Linde, Buzo & Gray 1980]
- This is essentially a version of k-means with a hierarchical initialization of class centers (codewords)
  - ✓ Start with 1 class
  - ✓ Split the codeword
  - ✓ Classify with 2 classes, and so on...



## LBG at work



## K-means and EM

### There is a close similarity between them. They both

- Compute a class membership function
  - ✓ k-means algorithm performs a *hard* assignment ( $u_{kj}$ )
  - ✓ EM performs a *soft* assignment ( $w_{kj}$ )
- Use it to weight the computation of the mean (EM: also dispersion & priors)

### The variable $u_{kj}$ plays a role similar to the one of $w_{kj}$ in the EM

- Some authors (e.g. Duda) denote them both as  $P(\omega_j | \mathbf{x}_k, \theta)$

### k-means may be derived as a limit case of EM for Gaussian mixtures with diagonal, isotropic covariance matrices [Bishop]

- The variance is fixed and the same for all classes.
- Consider the limit where the variance goes to 0
  - ✓ The EM weights  $w_{kj}$  go to 0 except for the closest class, for which it goes to 1.
  - ✓ One obtains the same expression for the mean

## Fuzzy C-means

[DUNN, 1973]

### In the K-means algorithm

- The class membership coefficient is binary

### In the fuzzy C-means, $u_{kj} \in [0,1]$ and $\sum_{j=1}^C u_{kj} = 1$

- i.e. one sample has a graded membership to all clusters

### One minimizes a heuristic global cost function

$$J_{\text{fuzzy}} = \sum_{j=1}^C \sum_{k=1}^N u_{kj}^b \|x_k - m_j\|^2$$


### The exponent, $b$ tunes the “fuzziness” of the algorithm, $b > 1$

## Fuzzy C-means: algorithm

- Same structure as the K-means
- Fix class centers and minimize w.r.t.  $u$

- Use a Lagrangian to incorporate unit sum condition
- Canceling derivatives:

$$u_{kj}^b = \frac{(1/d_{kj})^{1/(b-1)}}{\sum_{c=1}^C (1/d_{kc})^{1/(b-1)}}$$



- Fix  $u$  and update centers

- Canceling derivatives:

$$m_j = \frac{\sum_{k=1}^N u_{kj}^b x_k}{\sum_{k=1}^N u_{kj}^b}$$

- Convergence may be better than for the K-means...

## Density-based clustering


- Clusters: dense regions separated by regions of lower density

- DBSCAN (Ester et al, 1996)


- Neighborhood distance,  $\epsilon$
- Minimal neighbor number, minPoint
- Repeat
  - Find core point and expand cluster
- Until all points left are isolated (outliers)

- Pros / cons

- Number of clusters not needed
- Captures arbitrarily shaped clusters
- Handles noisy datasets / outliers
- Sensitive to dimensionality
- Difficulties with varying densities



DBSCAN cluster model  
[Schubert2020]



LIDAR data analysis using DBSCAN  
© Matlab 2020



## Fast clustering techniques

- Useful for initializing more sophisticated clustering algorithms
- Random selection of  $C$  samples
- Division of variables [Webb] :
  - Choose one variable (e.g. principal component), divide its range into  $C$  bins. Every sample is then classified according to the bin it falls into. Finally, cluster centers are chosen as the means of the  $C$  resulting groups.
- Incremental algorithms
  - E.g. Leader Algorithm [Hartigan, 1975]
    - The first sample becomes the leader of the first cluster
    - Samples are considered one by one, in turn until they are all labeled
      - If a sample is close enough to the leader of one class, then it is assigned to that class.
      - Otherwise, it becomes the leader of a new class
  - Low algorithmic complexity, the number of classes is not required, but a threshold must be tuned
  - Leader - follower : when a sample is assigned to a cluster, its leader moves in the direction of the newly classified point. Neural implementation : ART.

## Hierarchical clustering

- Some algorithms provide a hierarchical representation of data

- Such a tree-like representation is often called a dendrogram
- The root of the tree is the whole data set, the nodes represent sub-clusters
- The leaves are the individual samples



- The dendrogram also provides a measure of (dis-)similarity between groups

- ...which is not the case for Venn diagrams



## 1st family: splitting algorithms

- **Top-down or “divisive” algorithms**
  - Start with a single cluster that contains the whole data set
  - Repeat
    - ✓ Choose the “worst” cluster
    - ✓ Split it
  - Until the number of clusters is equal to the number of samples, N
- **Choosing the “worst” cluster ?**
  - E.g. according to the number of elements, to the variance
- **Splitting clusters ?**
  - E.g. perpendicularly to the direction of greatest variance...
- **This technique may be complicated and rather costly.**

⇒ Ascending techniques are often preferred.

## 2nd family: merging algorithms

- **Bottom-up or “agglomerative” algorithms**
  - Start with N singleton clusters
  - Repeat
    - ✓ Choose the 2 closest clusters
    - ✓ Merge them
  - Until the number of clusters is 1
- **Choosing the closest clusters ?**
  - According to the distance between groups
  - Some popular choices are:

$$d_{\min}(\omega_i, \omega_j) = \min_{\substack{x \in \omega_i \\ y \in \omega_j}} \|x - y\|$$

$$d_{\max}(\omega_i, \omega_j) = \max_{\substack{x \in \omega_i \\ y \in \omega_j}} \|x - y\|$$

$$d_{\text{average}}(\omega_i, \omega_j) = \frac{1}{n_i n_j} \sum_{x \in \omega_i} \sum_{y \in \omega_j} \|x - y\|$$

$$d_{\text{centroid}}(\omega_i, \omega_j) = \|\mu_i - \mu_j\|$$

## Main algorithms

- **Single linkage (or minimum algorithm)**
  - Uses the minimum distance between clusters: “min(min)”
  - Is rather versatile (e.g. it can handle concentric clusters)
  - Produces elongated clusters by linking clusters, without loops or circuits,
  - The result is dubbed *Minimal Spanning Tree*.
- **Complete linkage (or maximum algorithm)**
  - Uses the maximum distance between clusters: “min(max)”
  - Produces compact clusters
  - This produces graphs in which edges connect all of the nodes in a cluster.
  - Each resulting cluster is a *complete subgraph*.
- **Both distances are sensitive to outliers**
  - Using the average distance or the distance between centroids may be better.

## Single linkage: example

[Webb]

Consider the dissimilarity matrix

	1	2	3	4	5	6
1	0	4	13	24	12	8
2		0	10	22	11	10
3			0	7	3	9
4				0	6	18
5					0	8,5
6						0

1. We merge 5 and 3

	1	2	(3,5)	4	6
1	0	4	12	24	8
2		0	10	22	10
(3,5)			0	6	8,5
4				0	18
6					0

2. We merge 1 and 2


	(1,2)	(3,5)	4	6
(1,2)	0	10	22	8
(3,5)		0	6	8,5
4			0	18
6				0

3. We merge (3,5) and 4

	(1,2)	(3,4,5)	6
(1,2)	0	10	8
(3,4,5)		0	8,5
6			0

4. We merge (1,2) and 6

	(1,2,6)	(3,4,5)
(1,2,6)	0	8,5
(3,4,5)		0






## Complete linkage: exercise



- Same exercise with the following dissimilarity matrix

	1	2	3	4	5	6
1	0	4	13	24	12	8
2		0	10	22	11	10
3			0	7	3	9
4				0	6	18
5					0	8,5
6						0

Recall that we always merge the closest clusters (even if a max is used to evaluate the distances)



## Remarks

- Different distances between groups were proposed

- Each one has a particular effect on the results of the clustering algorithm
- Some distances may be related to the minimization of an objective function
- E.g. Ward's distance corresponds to the quadratic error criterion
  - ✓ Merging groups necessarily increases intra-cluster variance
  - ✓ Using Ward's distance minimizes this increase

$$d_{Ward}(\omega_i, \omega_j) = \frac{n_i n_j}{n_i + n_j} \|m_i - m_j\|^2$$

$$J_e = \sum_{i=1}^C \sum_{x \in \omega_i} \|x - m_i\|^2$$

- Choosing the best number of groups ?

- It is a recurring and difficult question... some solutions are
  - ✓ Monitoring the "lifetime" or persistence of clusters in the dendrogram
  - ✓ When an objective criterion is used (ex.  $J_e$ ), see how it evolves as the number of clusters changes
  - ✓ Use a self-similarity criterion
  - ✓ Use hypothesis testing [Duda]...

## Hierarchical Density Based Clustering (HDBSCAN)

- Introduced by Campello et al (2013, 2015)
- Hierarchical version of (a reinterpreted version of) DBSCAN

- Core distance

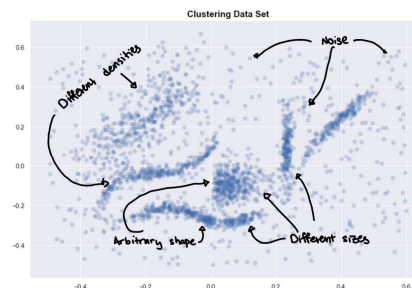
- $d_{core}$  = distance from a point to its k-th neighbor (same as DBSCAN)

- Hierarchical clustering

- Single linkage
- Mutual reachability distance  $\min\{d_{core}(x), d_{core}(y), d(x, y)\}$

- Flattening the hierarchy

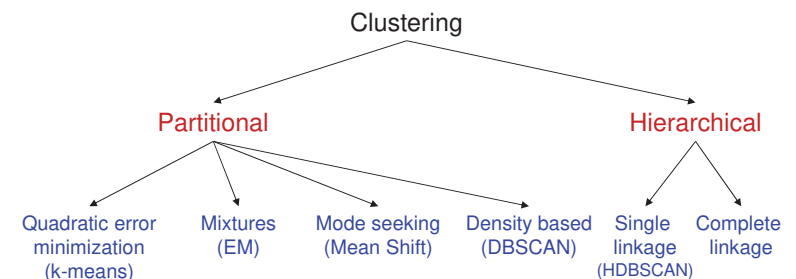
- Based on cluster persistence (cluster stability)



<https://towardsdatascience.com/a-gentle-introduction-to-hdbscan-and-density-based-clustering-5fd79329c1e8>

## Overview

(following [Jain99])



- Clustering methods:

- Deterministic vs. probabilistic
- Flat vs. hierarchical
- Monotheoretical vs. polytheoretical (use a single coordinate of x or the whole x)
- Hard vs. fuzzy
- Incremental or non-incremental

## Supplementary material

## A “good property” of the expectation [Lange2000]

### ▪ Jensen’s inequality for expectation:

$$E[h(x)] \geq h(E[x]) \text{ for } h \text{ convex}$$

### ▪ Taking $h = -\ln$ , one shows that for two pdf’s, $p$ and $q$

$$E_p[\ln(p)] - E_p[\ln(q)] \geq 0$$

- Demo:  $-\ln(\cdot)$  being convex,  $E_p[\ln(p)] - E_p[\ln(q)] = E_p\left[-\ln\left(\frac{q}{p}\right)\right] \geq -\ln\left(E_p\left[\frac{q}{p}\right]\right)$

$$-\ln E_p\left[\frac{q}{p}\right] = -\ln\left(\int \frac{q}{p} p\right) = -\ln\left(\int q\right) = \ln(1) = 0$$

### ▪ Now take $p = p_{Z|X, \Theta^n}(Z|X, \Theta^n)$ and $q = p_{Z|X, \Theta^{n+1}}(Z|X, \Theta^{n+1})$ ,

$$\text{i.e. } E_p[\ln(p)] = H(\Theta^n | \Theta^n) \text{ and } E_p[\ln(q)] = H(\Theta^{n+1} | \Theta^n)$$

$$H(\Theta^n | \Theta^n) - H(\Theta^{n+1} | \Theta^n) \geq 0$$



## M-step, Gaussian Mixture Model (scalar case)

### ▪ The function that is maximized during the M-step is

$$Q(\Theta | \Theta^n) = \sum_{j=1}^C \sum_{k=1}^N w_{kj}^n \log P(\omega_j) + \sum_{j=1}^C \sum_{k=1}^N w_{kj}^n \log p(x_k | \theta_j)$$

where  $\Theta = \{P(\omega_j), \theta_j\}$  are the parameters of the mixture, and  $\theta_j$  is the set of parameters of component  $\omega_j$ .

### ▪ In the following slides, we develop the calculation of the optimal parameters, for the Gaussian Mixture Model (GMM)

- As is well-known, the method consists in cancelling the derivative of the function with respect to (w.r.t.) the parameter in question.

## M-step, Gaussian Mixture Model (scalar case)

### ▪ Recall that if $x \sim \mathcal{N}(\mu, \sigma^2)$

$$\log p(x | \mu, \sigma^2) = -\frac{(x - \mu)^2}{2\sigma^2} - \log \sigma - \frac{1}{2} \log 2\pi$$

### ▪ Then, neglecting the additive constant, we have

$$Q(\Theta | \Theta^n) = \sum_{j=1}^C \sum_{k=1}^N w_{kj}^n \log P(\omega_j) - \sum_{j=1}^C \sum_{k=1}^N w_{kj}^n \frac{(x_k - \mu_j)^2}{2\sigma_j^2} - \sum_{j=1}^C \sum_{k=1}^N w_{kj}^n \log \sigma_j$$

### ▪ Maximizing $Q(\Theta | \Theta^n)$ w.r.t. $\mu_j$ only involves the 2<sup>nd</sup> term

$$2 \sum_{k=1}^N w_{kj}^n \frac{(x_k - \mu_j)}{2\sigma_j^2} = 0 \rightarrow \sum_{k=1}^N w_{kj}^n x_k = \mu_j \sum_{k=1}^N w_{kj}^n \Rightarrow \mu_j = \frac{\sum_{k=1}^N w_{kj}^n x_k}{\sum_{k=1}^N w_{kj}^n}$$

## M-step, Gaussian Mixture Model (scalar case)

$$Q(\theta|\theta^n) = \sum_{j=1}^C \sum_{k=1}^N w_{kj}^n \log P(\omega_j) - \sum_{j=1}^C \sum_{k=1}^N w_{kj}^n \frac{(x_k - \mu_j)^2}{2\sigma_j^2} - \sum_{j=1}^C \sum_{k=1}^N w_{kj}^n \log \sigma_j$$

- Maximizing  $Q(\theta|\theta^n)$  w.r.t.  $\sigma_j$  involves the last two terms

$$0 = \sum_{k=1}^N w_{kj}^n \frac{(x_k - \mu_j)^2}{\sigma_j^3} - \sum_{k=1}^N \frac{w_{kj}^n}{\sigma_j}$$

- Assuming  $\sigma_j$  to be non-null, we obtain

$$\sigma_j^2 = \frac{\sum_{k=1}^N w_{kj}^n (x_k - \mu_j)^2}{\sum_{k=1}^N w_{kj}^n}$$

## M-step, Gaussian Mixture Model (scalar case)

- For  $P(\omega_j)$  we use the constraint  $\sum_{j=1}^C P(\omega_j) = 1$  and maximize

$$\sum_{j=1}^C \sum_{k=1}^N w_{kj}^n \log P(\omega_j) + \lambda \left( 1 - \sum_{j=1}^C P(\omega_j) \right)$$

- Cancelling the derivative w.r.t.  $P(\omega_j)$ , we obtain

$$0 = \sum_{k=1}^N \frac{w_{kj}^n}{P(\omega_j)} - \lambda \rightarrow P(\omega_j) = \frac{1}{\lambda} \sum_{k=1}^N w_{kj}^n$$

- Summing over  $j$  and since  $\sum_{j=1}^C P(\omega_j) = 1$

$$\lambda = \sum_{j=1}^C \sum_{k=1}^N w_{kj}^n = N \Rightarrow P(\omega_j) = \frac{1}{N} \sum_{k=1}^N w_{kj}^n$$



## Solutions of exercises

### Unsupervised hierarchical clustering algorithms

## Complete linkage: exercise



Consider the dissimilarity matrix

	1	2	3	4	5	6
1	0	4	13	24	12	8
2		0	10	22	11	10
3			0	7	3	9
4				0	6	18
5					0	8,5
6						0

1. We merge 5 and 3

	1	2	(3,5)	4	6
1	0	4	13	24	8
2		0	11	22	10
(3,5)			0	7	9
4				0	18
6					0

2. We merge 1 and 2

	(1,2)	(3,5)	4	6
(1,2)	0	13	24	10
(3,5)		0	7	9
4			0	18
6				0

3. We merge (3,5) and 4

	(1,2)	(3,5,4)	6
(1,2)	0	24	10
(3,5,4)		0	18
6			0

4. We merge (1,2) and 6

	(1,2,6)	(3,5,4)
(1,2,6)	0	24
(3,5,4)		0

Recall that one always merge the **closest** clusters (even if a max is used to evaluate the distances)

