

## Support Vector Machines (SVM): an introduction

Linear discriminant with margin  
**Support Vector Machines**  
Kernels and non linear Support Vector Machines

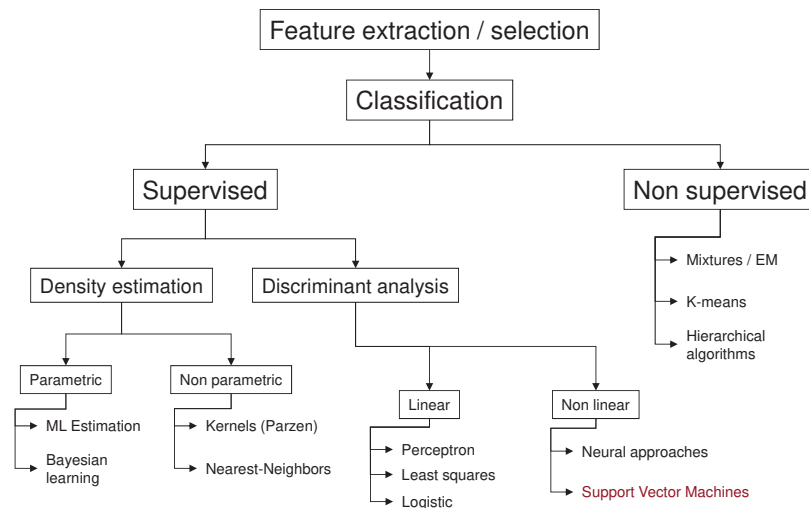
Special thanks to J.P. Tarel (Université Gustave Eiffel)



## Summary of previous episodes




- In this course, we studied several ways of devising linear discriminant functions
  - Parametric Gaussian models with identical variances in the Bayesian framework,
  - Direct estimation of the linear discriminant function by criterion minimization (Perceptron, Least Squares or Logistic regression, for example).
- We also studied 2 ways of generating discriminant functions in a neuro-mimetic (or neural) approach
  - Namely, multilayer Perceptrons and Radial basis Functions (or RBF's)
- We now introduce an alternative technique: **Support Vector Machines or SVMs (in French: “Séparateurs à Vaste Marge”)**.
  - Proposed in 1995, based on much older ideas (1963, 1964...)

## A hierarchy of methods



## Introduction – Basis ideas

- **Support Vector Machines (SVM)**
  - Essentially, a special case of linear discriminant.
  - Consider  $N$  samples drawn from two linearly separable classes.
  - Each sample  $\mathbf{x}_k$  has a label,  $t_k \in \{-1, 1\}$
  - Find a separating hyper-plane by a learning procedure
- **Multiple possible choices**
  - Solution: Margin maximization methods [Vapnik1963]
- **Classes may not be linearly separable**
  - Class overlap  
Solution: Slack variables [Cortes1995] / hinge loss
  - Nonlinearity  
Solution: Kernel methods [Aizerman1964]




## Maximum margin linear classifiers

- Perceptron with margin, where  $t_k \in \{-1, 1\}$

$$t_k(\mathbf{w}^T \mathbf{x}_k + w_0) \geq b$$

- Identical up to a scaling factor...
- Canonical Hyper-plane :  $|b| = 1$  for samples lying on margin hyper-planes  $H_1$  and  $H_2$
- The distance from these points to the separating surface (the margin) is  $1/\|\mathbf{w}\|$



- Maximizing the margin = constrained optimization problem

- Minimize  $\frac{1}{2} \|\mathbf{w}\|^2$  s.t.  $t_k(\mathbf{w}^T \mathbf{x}_k + w_0) \geq 1$ , for  $k = 1, \dots, N$

## Learning: the dual optimization problem

- Tool for constrained optimization: the Lagrangian

$$L_P = \frac{1}{2} \mathbf{w}^T \mathbf{w} - \sum_{k=1}^N \alpha_k (t_k (\mathbf{w}^T \mathbf{x}_k + w_0) - 1)$$

- Minimize w.r.t.  $\mathbf{w}$ ,  $w_0$  and maximize w.r.t.  $\alpha_k \geq 0$  (Lagrange multipliers)

$$\frac{\partial L_P}{\partial \mathbf{w}} = 0 \Rightarrow \mathbf{w} = \sum_{k=1}^N \alpha_k t_k \mathbf{x}_k \quad \frac{\partial L_P}{\partial w_0} = 0 \Rightarrow \sum_{k=1}^N \alpha_k t_k = 0$$

- Plugging these expressions into  $L_P$ , one obtains the following dual problem: maximize w.r.t.  $\alpha$

$$L_D = \sum_{k=1}^N \alpha_k - \frac{1}{2} \sum_{j=1}^N \sum_{k=1}^N \alpha_j \alpha_k t_j t_k (\mathbf{x}_j^T \mathbf{x}_k)$$

- Subject to the (simpler) constraints  $\alpha_k \geq 0$  and  $\sum_{k=1}^N \alpha_k t_k = 0$

## Learning: solving the dual problem

- Numerical methods are available to solve this quadratic programming problem and estimate  $\alpha^{opt}$ ,
  - e.g. Matlab's `quadprog` routine (optimization toolbox)
  - or Python's `cvxopt` or `quadprog` packages
- A constrained optimization problem of this sort satisfies the Kuhn, Karush and Tucker conditions

$$\alpha_k \geq 0$$

$$t_k (\mathbf{w}^T \mathbf{x}_k + w_0) - 1 \geq 0$$

$$\alpha_k (t_k (\mathbf{w}^T \mathbf{x}_k + w_0) - 1) = 0$$

- At least one factor must be zero. For most samples, it will be  $\alpha_k$

→ Points lying on  $H_1$  and  $H_2$  are the only ones for which  $\alpha_k = 0$  may not be satisfied. They are called Support Vectors (SV)

## Discriminant function

- Once the learning stage is performed, the discrimination rule may be simplified by keeping only the  $n \leq N$  SV: the solution is *sparse*
- Using the expression of  $\mathbf{w}$  as a function of  $\alpha^{opt}$ :

$$\mathbf{w} = \sum_{k=1}^N \alpha_k t_k \mathbf{x}_k \quad \Rightarrow \quad g(\mathbf{x}) = \sum_{k=1}^n \alpha_k^{opt} t_k (\mathbf{x}_k^T \mathbf{x}) + w_0^{opt}$$

- For  $w_0$ , one must come back to the primal problem,

- $t_i \cdot g(\mathbf{x}_i) = 1$  for all support vectors, i.e.  $t_i^2 \cdot g(\mathbf{x}_i) = g(\mathbf{x}_i) = t_i$  (since  $t_i^2 = 1$ )
- Replacing  $g(\mathbf{x}_i)$  by its expression and taking the mean

$$t_i - \sum_{k=1}^n \alpha_k^{opt} t_k (\mathbf{x}_k^T \mathbf{x}_i) = w_0^{opt} \quad \Rightarrow \quad w_0^{opt} = \frac{1}{n} \sum_{i=1}^n \left( t_i - \sum_{k=1}^n \alpha_k^{opt} t_k (\mathbf{x}_k^T \mathbf{x}_i) \right)$$

## Learning with overlapping class distributions

### The soft margin model [Cortes1995]


- Introduces *slack* variables,  $\xi_k \geq 0, k = 1, \dots, N$  (one per sample)
- Relaxes classification condition:  $t_k(\mathbf{w}^T \mathbf{x}_k + w_0) \geq 1 - \xi_k$ 
  - $\xi_k = 0$ : usual condition  $t_k(\mathbf{w}^T \mathbf{x}_k + w_0) \geq 1$
  - $0 < \xi_k \leq 1$ : allows margin violation,  $t_k(\mathbf{w}^T \mathbf{x}_k + w_0) \geq 0$
  - $\xi_k > 1$ : allows misclassification,  $t_k(\mathbf{w}^T \mathbf{x}_k + w_0) < 0$
- Penalizes relaxations:  $\sum_{k=1}^N \xi_k$

### The optimization problem becomes

- Minimize  $\frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{k=1}^N \xi_k$   
s.t.  $t_k(\mathbf{w}^T \mathbf{x}_k + w_0) \geq 1 - \xi_k$ ,  
 $\xi_k \geq 0, k = 1, \dots, N$

### Mathematical developments are similar

- Support Vectors satisfy:  $t_k(\mathbf{w}^T \mathbf{x}_k + w_0) = 1 - \xi_k$



Following [Bishop]

## Learning: solving the primal problem


### Dual problem: size $N$ (one Lagrange multiplier per sample)

### Primal: size $D + 1$ (one coefficient per dimension + bias)

### Reformulate primal criterion as unconstrained optimization

- Introducing  $g_k \stackrel{\text{def}}{=} \mathbf{w}^T \mathbf{x}_k + w_0$ , the constraint  $t_k g_k \geq 1 - \xi_k$  may be written as  $\xi_k \geq 1 - t_k g_k$  which, with  $\xi_k \geq 0$ , is equivalent to  $\xi_k^* = \max(0, 1 - t_k g_k)$
- The learning problem is formulated as the minimization, with respect to  $\mathbf{w}$ , of


$$\underbrace{\frac{1}{2} \|\mathbf{w}\|^2}_{\text{regularization}} + C \underbrace{\sum_{k=1}^N \max(0, 1 - t_k g_k)}_{\text{hinge loss } \mathcal{L}}$$



### Algorithm


- The hinge loss is not differentiable, but is convex and has a *sub-gradient*
- The primal criterion can be optimized using (stochastic) gradient descent

## Managing non linearly separable data



### Map $\mathbf{x}$ to $\Phi(\mathbf{x})$ to make data separable

- Use polar coordinates  
 $\mathbb{R}^2 \rightarrow \mathbb{R}^2$
- Or map data to higher dimension  
 $\mathbb{R}^2 \rightarrow \mathbb{R}^3$



## SVM in transformed feature space

### Map coordinates to transformed feature space

$$\Phi: \mathbb{R}^d \mapsto \mathbb{R}^D$$

$$\mathbf{x} \rightarrow \Phi(\mathbf{x})$$

### Learn a linear classifier for $\mathbf{w} \in \mathbb{R}^D$

$$g(\mathbf{x}) = \mathbf{w}^T \Phi(\mathbf{x}) + w_0$$

e.g. primal classifier formulation

$$\min_{\mathbf{w} \in \mathbb{R}^D} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{k=1}^N \max(0, 1 - t_k g(\mathbf{x}_k))$$

### If $D \gg d$ , there are many more parameters to learn for $\mathbf{w}$

→ Can we avoid this?

## Dual SVM in transformed feature space

- Expression of the discriminant function (dual)

$$g(\mathbf{x}) = \sum_{k=1}^N \alpha_k t_k \Phi(\mathbf{x}_k)^T \Phi(\mathbf{x}) + w_0$$

- Learning: the dual optimization problem

$$\begin{aligned} \text{Maximize } L_D &= \sum_{k=1}^N \alpha_k - \sum_{j=1}^N \sum_{k=1}^N \alpha_j \alpha_k t_j t_k \Phi(\mathbf{x}_j)^T \Phi(\mathbf{x}_k) \\ \text{subject to } \alpha_k &\geq 0 \text{ and } \sum_{k=1}^N \alpha_k t_k = 0 \end{aligned}$$

- Remarks

- Note that,  $\Phi$  appears in scalar products  $\Phi(\mathbf{x})^T \Phi(\mathbf{x}')$
- Once the scalar products are computed, only the  $N$  –dimensional vector  $\alpha$  needs to be learnt (instead of  $D + 1$  weights in the primal problem)

## The “Kernel Trick”

[Aizerman 1964]

- Write  $\Phi(\mathbf{x})^T \Phi(\mathbf{x}') = K(\mathbf{x}, \mathbf{x}')$ , where  $K$  is known as a **kernel**
- Expression of the discriminant function

$$g(\mathbf{x}) = \sum_{k=1}^N \alpha_k t_k K(\mathbf{x}_k, \mathbf{x}) + w_0$$

- Learning: the dual optimization problem

$$\begin{aligned} \text{Maximize } L_D &= \sum_{k=1}^N \alpha_k - \sum_{j=1}^N \sum_{k=1}^N \alpha_j \alpha_k t_j t_k K(\mathbf{x}_j, \mathbf{x}_k) \\ \text{subject to } \alpha_k &\geq 0 \text{ and } \sum_{k=1}^N \alpha_k t_k = 0 \end{aligned}$$

- The Kernel Trick

- Nor learning neither classification require explicitly computing the mapping,  $\Phi(\mathbf{x})$
- The mapping remains implicit: all that is needed is the kernel function,  $K(\mathbf{x}, \mathbf{x}')$
- Computing kernels is most often cheaper than computing scalar products in transformed feature space

## Example: quadratic kernels

- Consider the squared scalar product of  $\mathbf{x}, \mathbf{x}' \in \mathbb{R}^2$

$$K(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^T \mathbf{x}')^2 = \left( \begin{pmatrix} x_1 & x_2 \end{pmatrix} \begin{pmatrix} x'_1 \\ x'_2 \end{pmatrix} \right)^2 = (x_1 x'_1 + x_2 x'_2)^2$$

$$= (x_1 x'_1)^2 + 2x_1 x'_1 x_2 x'_2 + (x_2 x'_2)^2 = \begin{pmatrix} x_1^2 & \sqrt{2}x_1 x_2 & x_2^2 \end{pmatrix} \begin{pmatrix} x'^2_1 \\ \sqrt{2}x'_1 x'_2 \\ x'^2_2 \end{pmatrix}$$

$$\text{so } K(\mathbf{x}, \mathbf{x}') = \langle \Phi(\mathbf{x}), \Phi(\mathbf{x}') \rangle \text{ with } \Phi(\mathbf{x}) = \begin{pmatrix} x_1^2 & \sqrt{2}x_1 x_2 & x_2^2 \end{pmatrix}^T$$

- Note that, this decomposition is not unique. It also works for

- $\Phi(\mathbf{x}) = (1/\sqrt{2})(x_1^2 - x_2^2 \quad 2x_1 x_2 \quad x_1^2 + x_2^2)^T$  in  $\mathbb{R}^3$
- $\Phi(\mathbf{x}) = (x_1^2 \quad x_1 x_2 \quad x_1 x_2 \quad x_2^2)^T$  in  $\mathbb{R}^4$

- The simple polynomial kernel contains terms of degree 2 only

- $K(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^T \mathbf{x}' + c)^2$  with  $c > 0$  also contains linear and constant terms


## Dealing with the XOR with polynomial kernels

- Consider the kernel  $K(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^T \mathbf{x}' + 1)^2$

- After some manipulations (exercise):  
 $\Phi(\mathbf{x}) = (1, \sqrt{2}x_1, \sqrt{2}x_2, \sqrt{2}x_1 x_2, x_1^2, x_2^2)^T$ ,  $\mathcal{H} = \mathbb{R}^6$

- In the XOR case, (centered coordinates)

- $(-1, -1)$  and  $(1, 1)$  belong to class  $(t = 1)$ ,  $(1, -1)$  and  $(-1, 1)$  to class  $(t = -1)$
- “Manual” optimization leads to  $\alpha_k = 1/8 \quad \forall k$
- The 4 points are support vectors (unusual, due to the symmetry of the XOR)
- The resulting discriminant is:  $g(\mathbf{x}) = x_1 \cdot x_2$



Following [Duda]

## Mercer kernels

- Any function  $K$  that fulfills Mercer's condition may be used as a Kernel (no need to know the lifting function,  $\Phi$ , explicitly)

- If  $K(\mathbf{u}, \mathbf{v})$  is such that, for all square-integrable function  $f$  (i.e.  $\iint f(\mathbf{u})^2 d\mathbf{u} < \infty$ )  $\iint K(\mathbf{u}, \mathbf{v}) f(\mathbf{u}) f(\mathbf{v}) d\mathbf{u} d\mathbf{v} \geq 0$  then there exists a mapping:

$$\begin{aligned}\Phi: \mathcal{L} &\mapsto \mathcal{H} \\ \mathbf{x} &\rightarrow \Phi(\mathbf{x})\end{aligned}$$

with an expansion  $K(\mathbf{u}, \mathbf{v}) = \langle \Phi(\mathbf{u}), \Phi(\mathbf{v}) \rangle$  for all  $(\mathbf{u}, \mathbf{v}) \in \mathcal{L}^2$

- Kernel arithmetic** [Bishop, 2006]

- The following kernels:  $ck(\mathbf{x}, \mathbf{x}')$ ,  $f(\mathbf{x})k(\mathbf{x}, \mathbf{x}')f(\mathbf{x}')$ ,  $q(k(\mathbf{x}, \mathbf{x}'))$ ,  $\exp(k(\mathbf{x}, \mathbf{x}'))$ ,  $k_1(\mathbf{x}, \mathbf{x}') + k_2(\mathbf{x}, \mathbf{x}')$ ,  $k_1(\mathbf{x}, \mathbf{x}')k_2(\mathbf{x}, \mathbf{x}')$ , ... are valid kernels

where  $k(\mathbf{x}, \mathbf{x}')$ ,  $k_1(\mathbf{x}, \mathbf{x}')$ ,  $k_2(\mathbf{x}, \mathbf{x}')$  are valid kernels,  $c > 0$  is a constant,  $f$  is any function,  $q$  is a polynomial with nonnegative coefficients

## Well-known kernels

- Linear kernel**

$$K(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T \mathbf{x}'$$

- Polynomial kernels**

- Parameter  $M$ : the degree of the polynomial

$$K(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^T \mathbf{x}')^M$$

- Contains all monomials of order  $M$

$$K(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^T \mathbf{x}' + c)^M \text{ with } c > 0$$

- Contains all polynomial terms up to order  $M$

- Gaussian kernel**

- Parameter  $\sigma > 0$ : standard deviation


$$K(\mathbf{x}, \mathbf{x}') = \exp(-\|\mathbf{x} - \mathbf{x}'\|^2 / 2\sigma^2)$$

- No need to consider normalization constants in this context
- Dimension of transformed feature space = infinite  $\rightarrow$  cannot work in transformed space

## The Vapnik-Chervonenkis dimension

- The VC-dimension of a function is the maximum number of points that can be separated (*shattered*) using this function.


- E.g. oriented hyper-planes in  $\mathbb{R}^D$   
VC-dimension =  $D + 1$



- Polynomial kernels of order  $M$  have a VC-dimension of

$$\binom{d_{\mathcal{L}} + M - 1}{M} + 1$$

where  $d_{\mathcal{L}}$  is the dimension of  $\mathcal{L}$  (VC grows fast with  $M$ !)



From [Burges, 1998]

- The Gaussian kernel has an infinite VC-dimension


- Highly flexible, but beware of overfitting!

## Setting up hyper-parameters


- Care must be taken to the choice of the spread parameter of the Gaussian kernel

$$K(x_1, x_2) = \exp\left(-\frac{1}{2\sigma^2} \|x_1 - x_2\|^2\right)$$


- In practice: cross-validation, for example



$\sigma = 1.0$ , 10 SV



$\sigma = 0.01$ , 26 SV



$\sigma = 0.001$ , 27 SV


JP Tarel (UGE)

## Graphical interpretation


- Constructing the discriminant function  $g(\mathbf{x})$  amounts to positioning a kernel on each support vector, which defines a hyper-surface

$$\sum_{k=1}^N \alpha_k t_k K(\mathbf{x}_k, \mathbf{x})$$

- The decision boundary,  $g(\mathbf{x}) = 0$ , is obtained by “cutting” the hyper-surface at the altitude  $-w_0$




e.g. with Gaussian kernels



[Davy2003]

## Link with neural networks

- Dual SVMs may be seen as neural networks:



- Difference with 2-layer Perceptrons

- The number of hidden neurons is determined automatically by the support vectors  $x_1, \dots, x_n$ . Their output weights are the Lagrange multipliers,  $\alpha_k$ .

- Difference with RBF's

- The number of centers and their positions are automatically given by the learning stage of the SVM, as well as their weights (Lagrange multipliers) and threshold  $w_0$ .

## Take-away remarks

- Just as multilayer Perceptron's and RBF's,
  - SVM's can deal with non-linearly separable classes, thanks to the “kernel trick” and imperfectly separable classes, thanks to the “soft-margin model”
  - However, the number of parameters is much less, thanks to the maximal margin criterion
    - Spread parameter,  $\sigma$ , plus regularization parameter,  $C$ , for soft margins.
    - The other ones are automatically given by the learning procedure
- Learning = solve quadratic programming (optimization) problem
  - May be performed in a “reasonable” (polynomial) amount of time
  - Thanks to duality, the dimensionality does not matter...
  - Optimizing the primal problem is also feasible
- The kernel trick may be applied to any algorithm where the input vector enters only in the form of scalar products
  - Arises naturally in the dual form of SVMs
  - Perceptrons also may be *kernelized* (Aizerman, 1964), see Computer Exercise 16
  - Other *kernelized* algorithms include kernel PCA, Kernel Nearest-Neighbors classifier, Kernel Fisher discriminant, Kernel logistic regression...
- Extensions: 1-class SVM, multi-class SVM, regression SVM


## Ensemble learning

Ensemble learning  
Bagging - decision trees and random forests  
Boosting - AdaBoost

## Introduction: ensemble learning

### Methods that learn a target function by combining the predictions of a number of individual learners

- Models combinations are also called committees



### Why ensemble learning ?

- Decompose complex problems into multiple easier sub-problems
- Improve performance of individual (weak) learners
- No single model can solve all pattern recognition problems


### Possible strategies to generate classifiers

- Subsample the learning data set, then use subsets to train classifiers  
e.g. random sampling with replacement = bootstrap
- Train individual classifiers on different feature representations
- Use different training parameters (e.g.  $k$  in kNN) to generate classifiers

## Possible structures of ensemble classifiers


### Parallel (majority of approaches)

- Independent classifications
- Combination (average, majority vote, L-bit code word + error-correcting codes)



### Hierarchical (cascade)

- Sequential or tree-like combination
- First, inaccurate but fast classifiers
- Then, more accurate but more computationally intensive methods




## Motivating example

[Dietterich, 1997]

### Averaging can eliminate uncorrelated errors of classifiers


- e.g. combining 21 classifiers by majority vote
- Error rate of individual classifiers = 0.3 (30 %)
- Misclassification  $\Rightarrow$  11 classifiers in error over 21 (at least)
- Probability: 0.026 (2,6 %)



## Motivating example

### Plot of the ensemble error as a function of individual errors

- Base classifiers must be **better than random** (less than 50 % individual error)
- Another condition for success of the committee is **diversity** (i.e. uncorrelated errors)




### We will introduce two well-known strategies

- Train individual classifiers of the same kind on bootstrap samples of the training set and combine them by majority voting: **bagging**
- Train classifiers in sequence, adapting the training function to the performance of the previous model: **boosting**.


## Decision trees

(From [Criminisi2012])

### A general tree structure




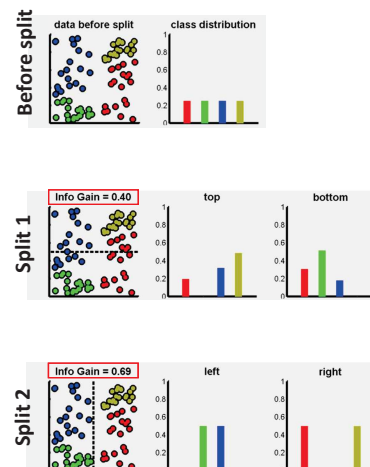
### A decision tree



## Building trees

(From [Criminisi2012])

- Each node corresponds to a **binary decision (threshold)**
  - Weak learner
- Decision based on an information - theoretic measure (entropy)
  - Promotes "peakness" or "purity"
- Random tree
  - Random selection of a subset of features at each node



## Random forests [Breiman2001]

### Bootstrap aggregating = bagging

### Learning


- Create **sub-samples** of learning data set by random selection with **replacement**
- Build one random tree per sub-sample → hence the name, **random forest**

### Decision

- Go down all trees
- Select class by **voting**

### Remarks

- Fast (parallel)
- Improves accuracy of unstable, uncorrelated classifiers



[Verikas2011]



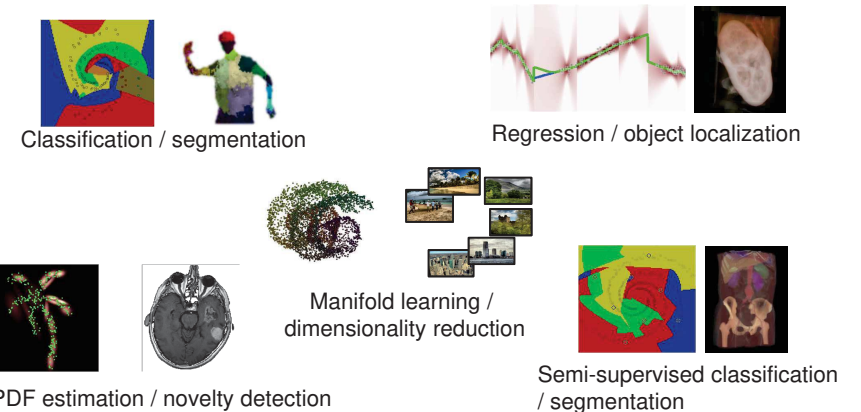
## Remarks

- **Out-of-bag (OOB) error**
  - Measures classification error on the data not used for training
  - Similar to (cross-)validation error
- **Variable importance**
  - Provides an ordered importance measure of the features used
  - Might be sensitive to the number of levels for categorical (discrete) features
- **Resources**
  - Decision Forests in Computer Vision and Medical Image Analysis. A. Criminisi and J. Shotton. Springer. 2013.
  - <http://research.microsoft.com/projects/decisionforests/>
  - Code
    - ✓ The Microsoft Research Cambridge Sherwood Software Library
    - ✓ Matlab (classification Toolbox), sklearn.ensemble.RandomForestClassifier
  - MOOCS (Nando de Freitas) : <http://www.youtube.com/watch?v=dCtJilEEgM>

## Decision forests

(From [Criminisi2012])

- The concept of classification and regression trees can be extended to other applications



## AdaBoost

[Freund&Schapire1995]

- Stands for adaptive boosting
- Combines weak classifiers
  - Misclassified samples are given more weight in successive learning step
- Intuitive example: learning kids how to recognize apples




"Apples are somewhat circular and somewhat red and possibly green"

From [Hsuan-Tien Lin,2009]

## AdaBoost

[Freund&Schapire1995]

- Example (weak learner = linear discriminant)



From [Antonio Torralba @MIT]

## AdaBoost algorithm

[Freund&Schapire1995]

- **Input: training examples**  $D=\{(x_n, t_n)\}_{n=1 \dots N}$
- **For**  $t = 1, 2 \dots T$ ,
  - Learn a simple rule  $h_t$  from emphasized training examples.
    - ✓ How? Choose a  $h_t \in H$  with minimum emphasized error.
  - Get the confidence  $\alpha_t$  of such rule
    - ✓ How? An  $h_t$  with lower error should get higher  $\alpha_t$
  - Emphasize the training examples that do not agree with  $h_t$ 
    - ✓ How? Maintain an emphasis value (weight)  $u_n$  per example

- **Output: combined function**

$$Y(x) = \text{sign}\left(\sum_{t=1}^T \alpha_t h_t(x)\right)$$

## Application to face detection

[Viola&Jones2001]

- **Widely-used method for face (and other objects) detection**
  - Detection, not recognition...i.e. binary classification
- **Very fast classification (real time), but slow training**
- **Viola&Jones method:**
  - Introduces Haar-like, fast-to-compute image features, that are well suited to face detection
  - Uses AdaBoost algorithm where each stage of the boosting process selects a single feature
  - Exploits a cascade of classifiers with increasing complexity and decreasing false alarm rate

## Features

- **Faces share specific properties**

- Eye region are darker than cheeks
- Nose bridge is brighter than eyes

== Domain-based knowledge


→ Can be captured by rectangle filters



[Viola&Jones2001]

- **Haar-like feature detector**

- Feature value =  $\sum$  (pixels in white rectangles) -  $\sum$  (pixels in black rectangles)
- Consider two- three- four-rectangle patterns
- All possible positions and sizes are considered
- ~160 000 possibilities in a 24x24 window



[Viola&Jones2001]

- **Fast computation using integral images**

## Integral images

- **Sum of pixels above and to the left of the current pixel**

- With  $i$ =original image,  $ii$  = integral image

$$ii(x, y) = \sum_{x' \leq x, y' \leq y} i(x', y')$$


- **Recursive implementation**

$$s(x, -1) = 0 \text{ and } ii(-1, y) = 0$$

$$s(x, y) = s(x, y-1) + i(x, y)$$

$$ii(x, y) = ii(x-1, y) + s(x, y)$$

- Where  $s$  is the cumulative row sum
- The integral image is computed only once, and in a single pass



[Viola&Jones2001]

- **Any rectangular sum (feature part) computed in constant time**


- e.g.  $1=A$ ,  $2=A+B$ ,  $3=A+C$ ,  $4=A+B+C+D \rightarrow D = 4+1-(2+3)$

## AdaBoost-feature selection

- On each round, many possibilities of weak classifiers
- 1 feature → 1 weak classifier
- At each stage of boosting
  - Given reweighted data from previous stage
  - Train all K (160,000) single-feature classifiers (threshold=decision stump)
  - Select the single best classifier (lowest weighted classification error)
  - Combine it with the other previously selected classifiers
  - Reweight the data
- Repeat until T classifiers selected
- Very computationally intensive
  - Learning K decision stumps T times
  - E.g., K = 160,000 and T = 1000
- Performs training + feature selection

## Cascade training

- A positive decision from the first classifier triggers the evaluation of the second one and so on...



[Viola&Jones2001]


- The classifiers are arranged by increasing complexity
  - 1 feature classifier → 100% detection rate, ~50% false positive rate
  - 5 features → 100% detection rate, ~40% FP rate (20% cumulative)
  - 20 features → 100% detection rate, 10% FP (2% cumulative)
- Viola&Jones: 38 layers in the cascade, 6060 features
  - 2, 10, 25, 25, 50, 50, 50 ... features.

## Results (1/2)

- Learning
  - 5000 faces x2 (vertical symmetry)
  - 10000 non-faces
- Learning takes “weeks”...
- Detection
  - 24x24 sliding window
  - Multi-scale: scale the detector, not the image (features can be computed at any scale)
  - In average, 10 features/window on test set
  - 10x faster than single-stage AdaBoost with 200 features
  - Due to analyzing window overlap, multiple nearby detections may occur → post-processing



[Viola&Jones2001]



## Results (2/2)



## Pros and cons

### Pros

- Extremely fast feature computation
- Scale and location invariant
- Very generic: can be trained to detect other facial features: nose, eyes, or body parts or many objects: cars, licence plates...
- Public implementations (and pre-trained cascades, also) available e.g. OpenCV, Matlab



### Cons

- Can hardly cope with face rotations (needs special training for profiles)
- Sensitive to lighting conditions

## Take-away

### Ensemble learning

- Multiple learners are trained to solve the same problem
- They try to build a set of diverse hypotheses and combine them for use

### Diversity

- Necessary to improve the performance of basic (weak) classifiers
- Can be introduced in different ways
  - ✓ Subsampling the training examples (e.g. bagging)
  - ✓ Manipulating the attributes, i.e. using subsets of a common feature representation
  - ✓ Using different training parameters
- => Injecting randomness into learning algorithms

### Strategies

- Bagging: use bootstrap (sampling with replacement), e.g. Random Forests
  - ✓ A "Swiss knife" for machine learning?
- Boosting: focusing on misclassified samples, e.g. AdaBoost

## Evaluation

### Methodology

### Confusion matrices


### Scoring metrics

### ROC / PR curves

## Evaluation: basis ideas

### Why?

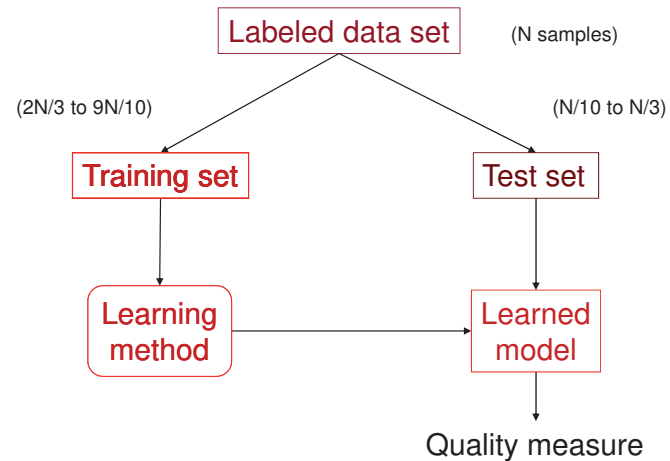
- Setting-up classifiers: compare methods/ optimize classifier parameters
- Assess performance (and limitations) of classifiers
  - ✓ No classifier is perfect !
  - ✓ Generalization capacities



### How?

- Needs a test data set different from training
  - ✓ Supervised: "ground-truth"
- What if not enough data? Generate sub-sets (cross-validation)
- Define a quality measure


## Test set definition (holdout method)



- Stratification: each class represented in both sets with same proportion


## Random resampling

- **Using a single training/test partition may be limited**
  - E.g. not enough data to make sufficiently large training and test sets
    - ✓ Large test set → more reliable estimate of performance
    - ✓ Large training set → more representative
- **Performance may be sensitive to the choice of the training set**
- **Solution: repeated random partitioning of the available data into training/test sets (repeated holdout) + average performance**



## Cross-validation (CV)

- Partition data into  $n$  subsamples
- Iteratively leave one subsample out for test, train on the rest
  - e.g.  $N=100$ , 5-fold cross validation



Iteration	Train on	Test on	Correct
1	$S_2 S_3 S_4 S_5$	$S_1$	16/20
2	$S_1 S_3 S_4 S_5$	$S_2$	15/20
3	$S_1 S_2 S_4 S_5$	$S_3$	17/20
4	$S_1 S_2 S_3 S_5$	$S_4$	14/20
5	$S_1 S_2 S_3 S_4$	$S_5$	18/20

Accuracy  
80 %

## Cross-validation


- Widely used approach for estimating test error
- Typical value:  $n=10$
- Leave-one-out Cross Validation (LOOCV):  $n=N$
- Stratified cross-validation: class proportions are maintained in each selected set
- Cross-validation vs. bootstrap: bootstrap tends to underestimate test error
  - Can be corrected: the “.632+” rule [Efron1997]

## Learning curves

- How does the performance of a learning method change as a function of the training-set size?

- Randomly select  $n$  instances from training set
- Learn model
- Evaluate model on test set
- One point on the curve

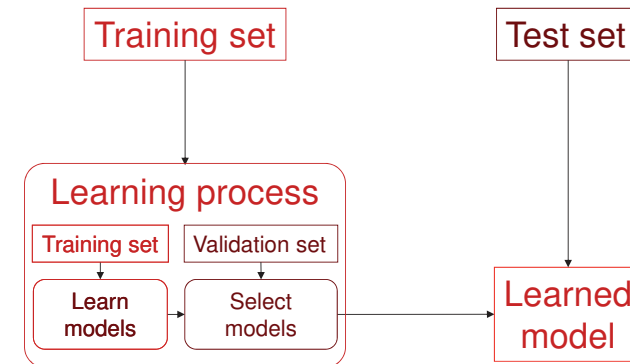
- Can be repeated several times (→ mean + error bars)



Source : <http://scikit-learn.org>

## Validation sets

- Used for tuning parameters during the learning process




- CV can be used inside learning step: nested cross-validation

## Validation score

- Plots the performance of the learned model as a function of a parameter, e.g. kernel width parameter for SVM
- Helps tuning the optimal value

- In general, three cases may be observed
  - Low training score, low validation score: underfitting
  - High scores: good performance
  - High training score, low validation score: overfitting



Source : <http://scikit-learn.org>

## Confusion matrix

- 2 classes
- $\omega_-$  negative  
 $\omega_+$  positive

		Prediction (result)		
		$\omega_-$	$\omega_+$	
Truth	$\omega_-$	TN – True negative	FP – False positive	N
	$\omega_+$	FN – False negative	TP – True positive	

Classification error	Accuracy	Recall	Precision
• $ERR = \frac{FP+FN}{Total}$	• $ACC = \frac{TP+TN}{Total}$	• $REC = \frac{TP}{FN+TP}$	• $PRE = \frac{TP}{FP+TP}$
		• What percentage of objects of interest are detected?	• What percentage of what we detect is correct?
		$F_1 = 2 \frac{PRE \times REC}{PRE + REC}$	

## Dealing with imbalanced classes

- Accuracy is not a good metric in case of imbalanced classes
  - E.g. breast cancer detection, 98 % negative, 2 % positive samples
  - Dummy classifier (no learning !) = always predict "negative"
    - 98% accuracy... but 0 % recall
- Example remedies to deal with imbalanced classes
  - Change performance metric (precision, recall, F1-score...)
  - Try to collect more data
  - Oversample minority class, or generate synthetic example (SMOTE)
  - Subsample majority class (at random, or by clustering)
  - Change the algorithm
    - ✓ Penalize wrong classification of the rare class (e.g. use weighted cross-entropy loss)
    - ✓ Tree based algorithms are considered to be less sensitive to the problem
    - ✓ Consider one-class classification + outlier detection

## Multi-class confusion matrix

- e.g. Optical Character (digits) Recognition with rejection class
- Observe confusions between, e.g. '4' and '9'

true class $i$	class $j$ predicted by a classifier										
	'0'	'1'	'2'	'3'	'4'	'5'	'6'	'7'	'8'	'9'	'R'
'0'	97	0	0	0	0	0	1	0	0	1	1
'1'	0	98	0	0	1	0	0	1	0	0	0
'2'	0	0	96	1	0	1	0	1	0	0	1
'3'	0	0	2	95	0	1	0	0	1	0	1
'4'	0	0	0	0	98	0	0	0	0	2	0
'5'	0	0	0	1	0	97	0	0	0	0	2
'6'	1	0	0	0	0	0	1	98	0	0	0
'7'	0	0	1	0	0	0	0	98	0	0	1
'8'	0	0	0	1	0	0	1	0	96	1	1
'9'	1	0	0	0	3	1	0	0	0	95	0

[Hlaváč]

## Scoring metrics for multiclass classification


- Accuracy : sum of diagonal elements /number of test samples
- Consider *one-versus-all* confusion matrices
  - $TN_c = \sum_{\mathbf{x} \notin \omega_c} (pred(\mathbf{x}_k) \neq \omega_c)$      $FP_c = \sum_{\mathbf{x} \notin \omega_c} (pred(\mathbf{x}_k) = \omega_c)$
  - $FN_c = \sum_{\mathbf{x} \in \omega_c} (pred(\mathbf{x}_k) \neq \omega_c)$      $TP_c = \sum_{\mathbf{x} \in \omega_c} (pred(\mathbf{x}_k) = \omega_c)$
- Micro-averaged metrics
  - e.g.  $PRE_{Micro} = \frac{TP_1 + \dots + TP_C}{TP_1 + \dots + TP_C + FP_1 + \dots + FP_C}$
  - Note that  $PRE_{Micro} = REC_{Micro} = ACC$
- Macro-averaged metrics (beware of class imbalance)
  - $PRE_{Macro} = \frac{PRE_1 + \dots + PRE_C}{C}$

## ROC curves

- ROC stand for Receiver Operational Characteristics
  - Terminology from RADAR detection
- Suited to binary classification (C=2)
  - Helps comparing methods, or tuning parameters

### ROC curve = parametric curve

- Most of the time, parameter = detection threshold
- For each parameter value,
  - ✓ Perform classification
  - ✓ Count TP / FA
  - ✓ Draw a point on the curve
- Example ROC curves
  - ✓ In green, ideal curve
  - ✓ In red, pure chance curve
  - ✓ In blue, 2 different classifiers
  - ✓ Performance metric: **Area under curve (AUC)**






## Precision-recall curves

### ■ Plots precision (relevance rate) vs. recall (True Positive Rate) as a function of a given parameter

- Most of the time, detection threshold
- Suited to the C=2 (dichotomy) case

### ■ Properties

- Not always monotonic
- End point (all samples selected)  
precision =  $P/(P+N)$
- Random classifier:  
constant precision  $P/(P+N)$
- The higher the curve, the better the classifier.
- Area under curve = **Average Precision**



## To summarize...

### ■ Evaluation is necessary for:

- Assessing the overall performance of a classifier (or a classification method)
  - ✓ Comparison with other classifiers
  - ✓ Tuning parameters
- Evaluating generalization capabilities

### ■ Take care to separate training from test data

- For tuning parameters, training data may be separated into training / validation

### ■ Cross validation is popular, especially when few data are available

### ■ Accuracy is not always the best evaluation criterion

- Sensitive to class imbalance
- Positive and negative errors may have different meanings
- Other criteria (precision, recall, specificity) may be envisioned

### ■ ROC curves and Precision-Recall curves are other common comparison tools for binary classification



## Conclusion

Hierarchy of methods  
Current trends in Pattern Recognition  
Towards *deep learning*

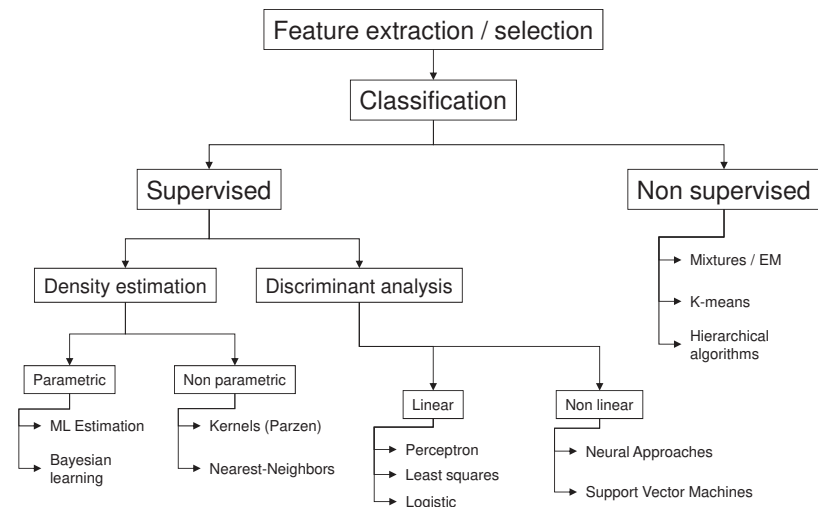
## Conclusion

- **Classification:** “The act of taking in raw data and making an action based on the “category” of the pattern” [Duda]
- **This implies two steps**
  - Machine learning (modeling classes)
  - Classification (making decisions)
- **We studied the main “traditional” statistical & connectionist pattern recognition approaches**
  - Feature extraction and selection
  - Bayesian decision and optimal classifiers
  - Supervised learning of probability density functions / priors
  - Supervised learning of linear / non-linear discriminant functions
  - Non-supervised learning (probabilistic, deterministic, hierarchical)


## Remark: classification and image processing

- **Both domains are closely related:**
  - Image processing is required to generate features from the images
    - ✓ Noise removal, contrast enhancement, primitive extraction, segmentation
  - Classification problems are many in I.P. & Computer Vision
    - ✓ e.g. thresholding, detection, segmentation
- **Classification for segmentation**
  - Pixel-wise decision
  - Accounting for pixel interaction: Markov Random Fields (OAT1 course)
    - ✓ Functional with 2 terms
      - Distance feature – class model
      - Regularization: label homogeneity on pixel's neighborhood.
  - Deformable regions (see course on Deformable Models)

## Hierarchy of methods




## Top 10 algorithms in data mining (IEEE, 2006)



- **KNN**
- **k-Means**
- Apriori
- **EM**
- **Naive Bayes**
- **Decision trees (and forests)**
  - C4.5,
  - CART
- **SVM**
- **AdaBoost**
- PageRank (Google)

...+ deep learning (2012: ImageNet)

## Trends in machine learning



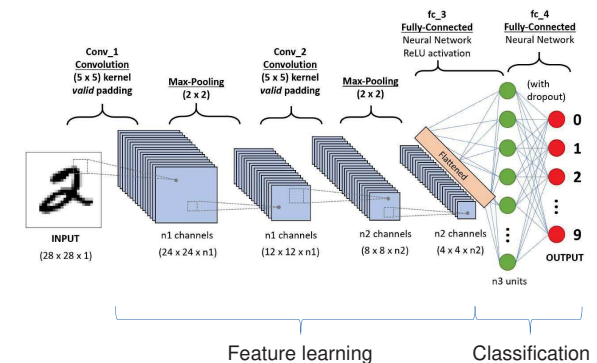
## 1985 - 2007: towards deep learning

- **NN are high-performance discrimination algorithms**
  - Boom in the 1990s, despite their "black box" side
  - But, still challenged by other algorithms during the early 2000's (SVM's, Random Forests, AdaBoost and cascades)
- **Convolutional neural networks (CNN)**
  - Neocognitron (Fukushima, 1980), LeNet (LeCun, 1989)
  - Reduce the number of connections and share weights
  - This mimics the behavior of *visual cortex*




## Convolutional neural networks (CNN)

- Input = 1 image, 1 neuron per image region (receptive field)
- Weight sharing, local connection → convolution
- Down-sampling → *pooling*
- Learning: back-propagation



## Feature learning vs. feature engineering



## The Deep Learning's boom

- **Other progresses in the late 2000's**
  - Bigger datasets
  - Faster computers (GPU's > 2007)
  - Better training algorithms and techniques (>2005)
- **The role of public competitions**
  - Public annotated dataset + competition + workshop
  - 2012: Deep Learning (AlexNet) wins the ImageNet contest
- **New research dissemination habits**
  - Publications (arXiv) and publicly available code (pre-trained NNs)
- **A flourishing research**
  - More and more applications, deeper networks, new architectures

*Thank you for your attention...*



## Credit

Many of the slides in this course  
were borrowed or adapted from Ricardo  
Gutierrez Osuna's Pattern Recognition course  
(Texas A&M University).



# Machine Learning and Pattern Recognition

## Bibliography

### Books and Theses

- [Biernat and Lutz, 2015] Biernat, E. and Lutz, M. (2015). *Data science : fondamentaux et études de cas*. Eyrolles, Paris.
- [Bishop, 2006] Bishop, C. M. (2006). *Pattern Recognition and Machine Learning (Information Science and Statistics)*. Springer-Verlag, Berlin, Heidelberg.
- [Criminisi and Shotton, 2013] Criminisi, A. and Shotton, J. (2013). *Decision forests for Computer Vision and Medical Image Analysis*. Springer, London.
- [Cristianini and Shawe-Taylor, 2000] Cristianini, N. and Shawe-Taylor, J. (2000). *An introduction to support vector machines (and other kernel-based learning methods)*. Cambridge University Press, Cambridge.
- [Duda et al., 2001] Duda, R. O., Hart, P. E., and Stork, D. G. (2001). *Pattern Classification*. John Wiley and Sons, Inc., New York, second edition.
- [Fukunaga, 1990] Fukunaga, K. (1990). *Introduction to Statistical Pattern Recognition*. Academic Press, New York, second edition.
- [Geron, 2019] Geron, A. (2019). *Machine Learning avec Scikit-Learn - Mise en œuvre et cas concrets*. Dunod, Malakoff.
- [Guyon et al., 2006] Guyon, I., Gunn, S., Nikravesh, M., and Zadeh, L. A., editors (2006). *Feature Extraction - Foundations and applications*. Springer, Berlin Heidelberg.
- [Hartigan, 1975] Hartigan, J. A. (1975). *Clustering Algorithms*. Wiley, New York.
- [Hastie et al., 2009] Hastie, T., Tibshirani, R., and Friedman, J. (2009). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Springer, second edition.
- [Hebb, 1949] Hebb, D. (1949). *The organization of behavior: a neuropsychological theory*. Wiley, New York.
- [Jurafsky and Martin, 2000] Jurafsky, D. and Martin, J. (2000). *Speech and Language Processing*. Pearson International Edition, first edition.

- [Minsky and Papert, 1969] Minsky, M. and Papert, S. (1969). *Perceptrons: An Introduction to Computational Geometry*. MIT Press, Cambridge, MA, USA.
- [Mitchell, 1997] Mitchell, T. (1997). *Machine Learning*. McGraw Hill.
- [Raschka and Mirjalili, 2017] Raschka, S. and Mirjalili, V. (2017). *Python Machine Learning*. Packt Publishing Ltd, Birmingham, second edition.
- [Salton and McGill, 1983] Salton, G. and McGill, M. (1983). *Introduction to modern information retrieval*. McGraw-Hill, New York.
- [Schalkoff, 1992] Schalkoff, R. (1992). *Pattern recognition: statistical, structural, and neural approaches*. Wiley, New York.
- [Theodoridis and Koutroumbas, 1999] Theodoridis, S. and Koutroumbas, K. (1999). *Pattern Recognition*. Academic Press, San Diego, first edition.
- [Webb and Copsey, 2011] Webb, A. R. and Copsey, K. D. (2011). *Statistical Pattern Recognition*. Wiley, Chichester, third edition.
- [Werbos, 1974] Werbos, P. (1974). *Beyond Regression: New Tools for Prediction and Analysis in the Behavioral Sciences*. PhD thesis, Harvard University.
- [Wu and Kumar, 2009] Wu, X. and Kumar, V. (2009). *The Top Ten Algorithms in Data Mining*. Chapman & Hall/CRC, Boca Raton, first edition.

### Papers, Book chapters and Technical reports

- [Aizerman et al., 1964] Aizerman, M., Braverman, E., and Rozonoer, L. (1964). Theoretical foundations of the potential function method in pattern recognition learning. *Automation and Remote Control*, 25(6):917–936. Translated from Russian.
- [Ball and Hall, 1965] Ball, G. and Hall, D. (1965). Isodata: a method of data analysis and pattern classification. Technical report, Stanford Research Institute, Menlo Park, United States. Prepared for Office of Naval Research. Information Sciences Branch.
- [Bilmes, 1998] Bilmes, J. A. (1998). A gentle tutorial on the EM algorithm and its application to parameter estimation for gaussian mixture and hidden markov models. Technical Report TR-97-021, University of Berkeley, Berkeley.

- [Boser et al., 1992] Boser, B., Guyon, I., and Vapnik, V. (1992). A training algorithm for optimal margin classifiers. In *Proceedings of the Fifth Annual Workshop on Computational Learning Theory*, pages 144–152.
- [Breiman, 2001] Breiman, L. (2001). Random forests. *Machine Learning*, 45(1):5–32.
- [Burges, 1998] Burges, C. J. C. (1998). A tutorial on support vector machines for pattern recognition. *Data Mining and Knowledge Discovery*, 2(2):121–167.
- [Comaniciu and Meer, 2002] Comaniciu, D. and Meer, P. (2002). Mean shift: A robust approach toward feature space analysis. *IEEE Trans. Pattern Analysis Machine Intell.*, 24(5):603–619.
- [Cortes and Vapnik, 1995] Cortes, C. and Vapnik, V. (1995). Support-vector networks. *Machine Learning*, 20(3):273–297.
- [Cover and Hart, 1967] Cover, T. and Hart, P. (1967). Nearest neighbor pattern classification. *IEEE Transactions on Information Theory*, 13(1):21–27.
- [Csurka et al., 2004] Csurka, G., Dance, C., Fan, L., Willamowski, J., and Bray, C. (2004). Visual categorization with bags of keypoints. *Workshop on Statistical Learning in Computer Vision, ECCV*, Vol. 1:1–22.
- [Cybenko, 1989] Cybenko, G. (1989). Approximation by superpositions of a sigmoidal function. *Mathematics of Control, Signals, and Systems*, 2:303–314.
- [Dempster et al., 1977] Dempster, A., Laird, N., and Rubin, D. (1977). Maximum likelihood from incomplete data via the EM algorithm. *Journal of the Royal Statistical Society. Series B (Methodological)*, 39(1):1–38.
- [Dietterich, 1997] Dietterich, T. (1997). Machine-learning research - four current directions. *AI Magazine*, 18(4):97–136.
- [Dunn, 1973] Dunn, J. C. (1973). A fuzzy relative of the isodata process and its use in detecting compact well-separated clusters. *Journal of Cybernetics*, 3(3):32–57.
- [Efron and Tibshirani, 1997] Efron, B. and Tibshirani, R. (1997). Improvements on cross-validation: The .632+ bootstrap method. *Journal of the American Statistical Association*, 92(438):548–560.
- [Ester et al., 1996] Ester, M., Kriegel, H., Sander, J., and Xu, X. (1996). A density-based algorithm for discovering clusters in large spatial databases with noise. In *Proceedings of the Second International Conference on Knowledge Discovery and Data Mining, KDD’96*, page 226–231. AAAI Press.
- [Fisher, 1936] Fisher, R. (1936). The use of multiple measurements in taxonomic problems. *Annals of Eugenics*, 7(2):179–188.
- [Fix and J.L. Hodges, 1951] Fix, E. and J.L. Hodges, J. (1951). Discriminatory analysis, nonparametric discrimination. Report 4, USAF School of Aviation Medicine, Randolph Field, Texas, USA. Project 21-49-004, Contract AF41(128)-31.
- [Freund and Schapire, 1997] Freund, Y. and Schapire, R. (1997). A decision-theoretic generalization of on-line learning and an application to boosting. *Journal of Computer and System Sciences*, 55(1):119–139.
- [Jain et al., 2000] Jain, A. K., Duin, R. P. W., and Mao, J. (2000). Statistical pattern recognition: a review. *IEEE Trans. Pattern Analysis Machine Intell.*, 22(1):4–37.
- [Jain et al., 1999] Jain, A. K., Murty, M. N., and Flynn, P. J. (1999). Data clustering: a review. *ACM Computing Surveys*, 31(3):264–323.
- [Lange et al., 2000] Lange, K., Hunter, D., and Yang, I. (2000). Optimization transfer using surrogate objective functions. *Journal of Computational and Graphical Statistics*, 9(1):1–20.
- [Le Cun, 1986] Le Cun, Y. (1986). Learning processes in an asymmetric threshold network. In Bienenstock, E., Fogelman-Soulié, F., and Weisbuch, G., editors, *Disordered systems and biological organization*, pages 233–240, Les Houches, France. Springer-Verlag.
- [Lee et al., 2011] Lee, H., Grosse, R., Ranganath, R., and Ng, A. Y. (2011). Unsupervised learning of hierarchical representations with convolutional deep belief networks. *Commun. ACM*, 54(10):95–103.
- [Linde et al., 1980] Linde, Y., Buzo, A., and Gray, R. (1980). An algorithm for vector quantizer design. *IEEE Transactions on Communications*, 28(1):84–95.
- [Lloyd, 1982] Lloyd, S. (1982). Least squares quantization in PCM. *IEEE Transactions on Information Theory*, 28(2):129–137.

- [MacQueen, 1967] MacQueen, J. (1967). Some methods for classification and analysis of multivariate observations. In *Proceedings of the Fifth Berkeley Symposium on Mathematical Statistics and Probability, Volume 1: Statistics*, pages 281–297. University of California Press.
- [McCulloch and Pitts, 1943] McCulloch, W. and Pitts, W. (1943). A logical calculus of the ideas immanent in nervous activity. *Bulletin of Mathematical Biophysics*, 5:115–133.
- [Moghaddam and Pentland, 2000] Moghaddam, B. and Pentland, A. (2000). Probabilistic visual learning for object representation. *IEEE Trans. Pattern Analysis Machine Intell.*, 19(7):696–710.
- [Parzen, 1962] Parzen, E. (1962). On estimation of a probability density function and mode. *Annals of Mathematical Statistics*, 33(3):1065–1076.
- [Powell, 1987] Powell, J. (1987). Radial basis functions for multivariable interpolation: a review. In Mason, J. and Cox, M., editors, *Algorithms for Approximation*, pages 143–167. Oxford University Press, Oxford.
- [Rosenblatt, 1957] Rosenblatt, F. (1957). The Perceptron - a perceiving and recognizing automaton. Technical Report 85-460-1, Cornell aeronautical laboratory, Inc., Buffalo.
- [Roweis, 1997] Roweis, S. (1997). EM algorithms for PCA and SPCA. In *NIPS*, pages 626–632.
- [Rumelhart et al., 1986] Rumelhart, D., Hinton, G., and Williams, R. (1986). Learning representations by back-propagating errors. *Nature*, 323(6088):533–536.
- [Tenenbaum et al., 2000] Tenenbaum, J., de Silva, V., and Langford, J. (2000). A global geometric framework for nonlinear dimensionality reduction. *Science*, 290(5500):2319–2323.
- [Tipping and Bishop, 1999] Tipping, M. E. and Bishop, C. M. (1999). Probabilistic principal component analysis. *Journal of the Royal Statistical Society. Series B (Statistical Methodology)*, 61(3):611–622.
- [Trunk, 1979] Trunk, G. V. (1979). A problem of dimensionality: A simple example. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 1(3):306–307.
- [Vapnik, 1963] Vapnik, V. (1963). Pattern recognition using generalized portrait method. *Automation and Remote Control*, 24:774–780.
- [Verikas et al., 2011] Verikas, A., Gelzinis, A., and Bacauskiene, M. (2011). Mining data with random forests: A survey and results of new tests. *Pattern Recognition*, 44(2):330–349.
- [Viola and Jones, 2001] Viola, P. and Jones, M. (2001). Rapid object detection using a boosted cascade of simple features. In *Proceedings of the 2001 IEEE Computer Society Conference on Computer Vision and Pattern Recognition. CVPR 2001*, volume 1, pages I–511–I–518.
- [Widrow and Hoff, 1960] Widrow, B. and Hoff, M. (1960). Adaptive switching circuits. *IRE WESCON Convention Record*, 4:96–104.
- [Wu et al., 2008] Wu, X., Kumar, V., Quinlan, J., Ghosh, J., Yang, Q., Motoda, H., McLachlan, G., Ng, A., Liu, B., Yu, P., Zhou, Z. H., Steinbach, M., Hand, D., and Steinberg, D. (2008). Top 10 algorithms in data mining. *Knowledge and Information Systems*, 14(1):1–37.

